



Swiss Federal Institute of Technology Zurich

Seminar for
Statistics

Department of Mathematics

Bachelor Thesis

Winter 2019

Nicolas Trutmann

Comparison of EM-algorithm and MLE using
Cholesky decomposition

Submission Date: placeholder

Advisor: placeholder

Abstract

The intent of this work is to compare The EM algorithm to a MLE approach in the case of multivariate normal mixture models using the Cholesky decomposition. The EM algorithm is widely used in statistics and is proven to converge, however in pathological cases convergence slows down considerably.

methods(not done)

results(not done)

Contents

1	Introduction to normal mixture models	1
1.1	Definitions	1
1.2	The EM-Algorithm in Sketch	2
1.3	Choice of Notation	3
1.4	Models of Covariance Matrices	4
1.5	Problems of the EM-algorithm	6
1.6	Alternative Option	7
2	The <code>norMmix</code> Package	9
2.1	Introduction to the Package	9
2.1.1	<code>norMmixMLE</code>	10
2.2	On The Development of <code>norMmix</code>	12
2.3	Demonstration	13
3	Comparing Algorithms	15
3.1	Time Analysis	17
3.2	Behaviour in <code>n</code>	18
3.3	Behaviour in <code>p</code>	18
3.4	Difficult Mixtures	19
3.5	Nonnormal Mixtures	21
4	Discussion	27
	Bibliography	28
A	R Code	31
A.1	<code>llnorMmix</code>	31
A.2	Example Simulation Script	36
B	Further Plots	37
B.1	Chapter 3	37

List of Figures

1.1	Parameters of <code>MW.nm9</code>	7
1.2	True and Estimated density	7
1.3	20 EM steps	7
1.4	200 EM steps	7
1.5	Log-likelihood Plotted against Iteration Count for the Example in 1.5 . . .	8
2.1	Demonstration of the MW Objects	13
2.2	The Correct Mixture	14
2.3	The Fitted Mixture	14
3.1	Example of Comparison Plot	16
3.2	Log-log Plot of System Time against Parameter Length	18

List of Tables

1.1	Table of Parameters of the Covariance Matrices	5
1.2	Full Table of Parameters	6
2.1	Translation Table: Mathematical Notation to R Code	9

Chapter 1

Introduction to normal mixture models

1.1 Definitions

A good and thorough introductory book is the work of [McLachlan and Peel \(2000\)](#) and the reader is encouraged to study it to learn in depth about normal mixtures and clustering. We will here give a short overview of normal mixtures to fix notation and nomenclature. The motivating idea behind mixture models is, that in real world examples a sample might be suspected to arise from more than one population or be more simply modelled by several overlayed distributions. The example of this, that is generally considered to be the first of this kind, is the one by Karl Pearson, who fitted two normal distributions with different means and variances. In his book, [Pearson and Henrici \(1896\)](#)[Section 4.d.; page 266], Pearson analyzed measurements of forehead to body length of crabs sampled from the bay of Naples. His mixture model-based approach suggested, that the crabs were evolving into two new subspecies. This is a historically important example, because it presents statistical evidence of evolution in process. Mixture models have been used since, but research took off after the availability of computing power made computational research possible

While the theory of mixture models holds for a much broader class of distributions, we restrict ourselves here to the case of normal distributions, because this restriction fits more comfortably into the scope of this work and because normal distributions allow for a parsimonious parametrization, that is of interest to study.

This parametrization is the LDL^T decomposition, which allows a very simple parametrization and a straightforward connection between degrees of freedom and necessarily generated numerical values. This will be explained further in section 1.4.

But before we delve deeper into the topic of this research, we first define the concept of a normal mixture model:

Let $\mu \in \mathbb{R}^p$, $\Sigma \in \mathbb{R}^{p \times p}$ be symmetric positive definite and $\phi(-; \mu, \Sigma)$ be the normal

29 distribution with mean μ and covariance matrix Σ with density function:

$$\phi(\mathbf{x}; \mu, \Sigma) = \frac{\exp(-\frac{1}{2}(\mathbf{x} - \mu)\Sigma^{-\frac{1}{2}}(\mathbf{x} - \mu)^\top)}{\sqrt{(2\pi)^k \det \Sigma}} \quad (1.1.0.1)$$

30 for $\mathbf{x} \in \mathbb{R}^p$. Since we are studying mixture models, we will need several overlapping of
 31 normal distributions, of differing means and covariance. Therefore, we choose notation
 32 allowing us to refer to the components in shorthand. Let us assume we have $K \in \mathbb{N}$
 33 normal distributions with means and covariance μ_k, Σ_k , $k \in \{1, \dots, K\}$, then we fix:

$$\phi_k(\mathbf{x}) := \phi(\mathbf{x}; \mu_k, \Sigma_k) \quad (1.1.0.2)$$

34 And going forward, we will refer to components by the subscript k .

35 **Definition 1.1.0.1.** Suppose we have a random sample $\mathbf{Y}_1, \dots, \mathbf{Y}_n$, where \mathbf{Y}_i is a p -
 36 dimensional random vector with probability density function $\mathbf{Y}_i \sim f(\mathbf{y}_i)$ on \mathbb{R}^p .

37 We assume that the density $f(\mathbf{y}_i)$ of \mathbf{Y}_i can be written in the form:

$$f(\mathbf{y}_i) = \sum_{k=1}^K \pi_k \phi_k(\mathbf{y}_i) \quad (1.1.0.3)$$

38 The ϕ_k are normal distributions and are called the mixture components with parameters
 39 μ_k and Σ_k as described above (1.1.0.2). The π_k are called the component densities of the
 40 mixture and are constrained by the rules $\pi_k > 0$ and $\sum_k \pi_k = 1$.

41 For 'large' datasets there are more parsimonious parametrizations, that reduce computa-
 42 tion time. These, for example, assume that all components have the same covariance, or
 43 have certain restrictions placed on them. We will give a detailed description of the models
 44 assumed in this thesis in section 1.4.

45 1.2 The EM-Algorithm in Sketch

46 With this definition, we immediately face the problem of how to fit these mixture com-
 47 ponents to given data. A popular algorithm to solve this problem is the **Expectation-**
 48 **Maximization** algorithm, abbreviated as EM-algorithm.

49 We give here a sketch of the EM-algorithm in the case of all normal mixture components.
 50 This roughly follows the content in [McLachlan and Peel \(2000\)](#). For a more thorough
 51 treatment of the matter see chapter 3.

52 Suppose we have a p -dimensional dataset of n samples x_1, \dots, x_n , onto which we would
 53 like to fit a K component normal mixture with mixture components ϕ_k , $k \in 1, \dots, n$.

54 For the EM-algorithm further parameters are introduced. These are denoted $\tau_j(\mathbf{y}_i)$ and
 55 they represent the posterior probabilities that observation i is a member of component j .

56 The EM-algorithm is a two step, iterative process consisting of an 'e'-step and an 'm'-step.
 57 In the e-step the expectation of component membership is updated.

$$\tau_j(\mathbf{y}_i; \Psi) = \phi_j(\mathbf{y}_i) / \sum_{k=1}^K \phi_k(\mathbf{y}_i) \quad (1.2.0.1)$$

and in the m-step given the component membership information we update the component means and covariances by weighted versions of the usual estimators.

$$\boldsymbol{\mu}_j = \sum_{i=1}^n \tau_j(\mathbf{y}_i) \mathbf{y}_i / \sum_{j=1}^n \tau_j(\mathbf{y}_i) \quad (1.2.0.2)$$

$$\boldsymbol{\Sigma}_j = \sum_{i=1}^n \tau_j(\mathbf{y}_i) (\mathbf{y}_i - \boldsymbol{\mu}_j)(\mathbf{y}_i - \boldsymbol{\mu}_j)^\top / \sum_{i=1}^n \tau_j(\mathbf{y}_i) \quad (1.2.0.3)$$

There remains to be stated how to start the algorithm. Since both steps of the algorithm depend on data from the other, the EM-algorithm needs some form of initialization step. Most popular implementations use some form of pre clustering and use the EM-algorithm as subsequent tools to fit the data. The R-package `mclust` for example uses hierarchical agglomerative clustering [Scrucca, Fop, Murphy, and Raftery \(2016\)](#).

1.3 Choice of Notation

The classification of models in this paper relies heavily on the work of [Celeux and Govaert \(1995\)](#), however, out of necessity for clarity, we break with their notation. So as to not confuse the reader we describe here in depth the differences in notation between [Celeux and Govaert \(1995\)](#) and ours.

The basis of classification in [Celeux and Govaert \(1995\)](#) is the decomposition of a symmetric matrix into an orthogonal and a diagonal component. A symmetric positive definite matrix Σ can be decomposed as follows

$$\Sigma = \lambda \mathbf{D} \mathbf{A} \mathbf{D}^\top \quad (1.3.0.1)$$

with \mathbf{D} an orthogonal matrix and \mathbf{A} a diagonal matrix and $\lambda = \sqrt[p]{\det(\Sigma)}$ the p -th root of the determinant of Σ .

This decomposition has an appealing geometric interpretation, with \mathbf{D} as the *orientation* of the distribution, \mathbf{A} the *shape*, and λ the *volume*. The problem of notation comes from standard conventions in linear algebra, where the letters A and D are usually occupied by arbitrary and diagonal matrices respectively. Furthermore, we intend to apply a variant of the Cholesky decomposition to Σ , the $\alpha \mathbf{L} \mathbf{D} \mathbf{L}^\top$ decomposition. This obviously raises some conflicts in notation.

Therefore we, from here on, when referring to the decomposition as described by [Celeux and Govaert \(1995\)](#), will use the following modification of notation:

$$\mathbf{D} \mapsto \mathbf{Q} \quad (1.3.0.2)$$

$$\mathbf{A} \mapsto \boldsymbol{\Lambda} \quad (1.3.0.3)$$

$$\lambda \mapsto \alpha \quad (1.3.0.4)$$

$$\Sigma = \lambda \mathbf{D} \mathbf{A} \mathbf{D}^\top = \alpha \mathbf{Q} \boldsymbol{\Lambda} \mathbf{Q}^\top \quad (1.3.0.5)$$

These were chosen according to general conventions of linear algebra. \mathbf{Q} is usually chosen for orthonormal matrices; $\boldsymbol{\Lambda}$ is often a choice for diagonal matrices of eigenvectors and α was somewhat arbitrarily chosen.

1.4 Models of Covariance Matrices

As mentioned above, there are ways to constrain the covariance matrices for computational reasons. There are instances where the resulting loss of information is seen as acceptable, for example if, through consideration of the data, that a simplified model is acceptable. Another is if the sheer size of the data makes application of full generality impossible.

- We restrict the complexity of the decomposition components
- We restrict the variability of the mixture components

Let us look at the first case. We take the decomposition of a covariance matrix as $\Sigma = \alpha \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^\top$. Of these, we can simplify the structure of \mathbf{Q} and $\mathbf{\Lambda}$, by replacing them with the identity. If we set $\mathbf{Q} = \text{Id}$, we lose the freedom of orientation and if we set $\mathbf{\Lambda} = \text{Id}$ we restrict ourselves to spherical distributions.

of course, we cannot restrict $\boldsymbol{\lambda}$ while letting \mathbf{q} free, since

$$\mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^\top = \mathbf{Q} \text{Id} \mathbf{Q}^\top = \text{Id} \quad (1.4.0.1)$$

The second restriction simply means we hold the decomposition fixed throughout all covariance matrices. There is however an issue with the cholesky decomposition. For 10 out of 14 cases as defined by [Celeux and Govaert \(1995\)](#), there exists a canonical translation of decompositions. The 6 diagonal cases need no translation; the eigen and cholesky decomposition are equal to identity. For the non-diagonal cases note that for a given sym. pos. def. matrix Σ we have decompositions:

$$\Sigma = \alpha \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^\top \quad \Sigma = \alpha \mathbf{L} \mathbf{D} \mathbf{L}^\top \quad (1.4.0.2)$$

Since in both cases the enclosing matrices \mathbf{Q} and \mathbf{L} have determinant 1 the determinant of Σ falls entirely on α . therefore α , in these particular decompositions, is equal for both. [Celeux and Govaert \(1995\)](#) vary σ by either varying or holding fixed the volume (α/α_k), shape ($\mathbf{\Lambda}/\mathbf{\Lambda}_k$) and orientation (\mathbf{Q}/\mathbf{Q}_k). These 3 times 2 cases would yield the 8 out of 14 cases of non-diagonal cases. However there is no canonical transform for either variable orientation and fixed shape or fixed orientation and variable shape. The reason for this is that in the $\mathbf{L} \mathbf{D} \mathbf{L}^\top$ decomposition the lower diagonal matrix \mathbf{L} holds some of the shape of the matrix, which in the eigendecomposition is in the $\mathbf{\Lambda}$ matrix. In fact, \mathbf{L} is orthogonal if and only if $\mathbf{L} = \text{Id}_{n \times n}$. Therefore we can only decompose matrices where either both or neither shape and orientation vary. See table 1.1.

While we could in theory construct the cases $\mathbf{L} \mathbf{D}_k \mathbf{L}^\top$ and $\mathbf{L}_k \mathbf{D} \mathbf{L}_k^\top$, however they do not correspond to the desired geometric intent behind the differentiation of models and are therefore not included.

Model	Σ_k C&G	volume	shape	orientation	parameters	LDL^\top as in C&G	parameters	count
EII	αI	equal	equal	-	α	as in C&G		1
VII	$\alpha_k I$	var.	equal	-	α_k			K
EEl	αA	equal	equal	coord. axes	α, λ_i			$1 + (p-1)$
VEI	$\alpha_k A$	var.	equal	coord. axes	α_k, λ_i			$K + (p-1)$
EVI	αA_k	equal	var.	coord. axes	$\alpha, \lambda_{i,k}$			$1 + K(p-1)$
VVI	$\alpha_k A_k$	var.	var.	coord. axes	$\alpha_k, \lambda_{i,k}$			$K + K(p-1)$
EEE	$\alpha Q A Q^\top$	equal	equal	equal	$\alpha, \lambda_i, q_{i,j}$	αLDL^\top	$\lambda, d_i, l_{i,j}$	$1 + (p-1) + \frac{p(p-1)}{2}$
EVE	$\alpha Q A_k Q^\top$	equal	var.	equal	$\alpha, \lambda_{i,k}, q_{i,j}$	doesn't exist		$1 + K(p-1) + \frac{p(p-1)}{2}$
VEE	$\alpha_k Q A Q^\top$	var.	equal	equal	$\alpha_k, \lambda_i, q_{i,j}$	$\alpha_k LDL^\top$	$\lambda_k, d_i, l_{i,j}$	$K + (p-1) + \frac{p(p-1)}{2}$
VVE	$\alpha_k Q A_k Q^\top$	var.	var.	equal	$\alpha_k, \lambda_{i,k}, q_{i,j}$			$K + K(p-1) + \frac{p(p-1)}{2}$
EEV	$\alpha Q_k A Q_k^\top$	equal	equal	var.	$\alpha, \lambda_i, q_{i,j,k}$	don't exist		$1 + (p-1) + K \frac{p(p-1)}{2}$
VEV	$\alpha_k Q_k A Q_k^\top$	var.	equal	var.	$\alpha_k, \lambda_i, q_{i,j,k}$			$K + (p-1) + K \frac{p(p-1)}{2}$
EVV	$\alpha Q_k A_k Q_k^\top$	equal	var.	var.	$\alpha, \lambda_i, q_{i,j,k}$	$\alpha L_k D_k L_k^\top$	$\lambda, d_{i,k}, l_{i,j,k} \quad j > i$	$1 + K(p-1) + K \frac{p(p-1)}{2}$
VVV	$\alpha_k Q_k A_k Q_k^\top$	var.	var.	var.	$\alpha_k, \lambda_i, q_{i,j,k}$	$\alpha_k L_k D_k L_k^\top$	$\lambda_k, d_{i,k}, l_{i,j,k} \quad j > i$	$K + K(p-1) + K \frac{p(p-1)}{2}$

Table 1.1: Table of Parameters of the Covariance Matrices

Σ model	μ, π	Σ	reduced	$\mathcal{O}()$
EII	$K - 1 + pK$	1	$Kp + K$	$Kp + K$
VII	$K - 1 + pK$	K	$Kp + 2K - 1$	$Kp + K$
EEI	$K - 1 + pK$	$1 + (p - 1)$	$Kp + p + K - 1$	$Kp + p + K$
VEI	$K - 1 + pK$	$K + (p - 1)$	$Kp + p + 2K - 2$	$Kp + p + K$
EVI	$K - 1 + pK$	$1 + K(p - 1)$	$2Kp$	Kp
VVI	$K - 1 + pK$	$K + K(p - 1)$	$2Kp + K - 1$	Kp
EEE	$K - 1 + pK$	$1 + (p - 1) + \frac{p(p-1)}{2}$	$\frac{(p+2)(p-1)}{2} + Kp + K$	$p^2 + Kp + K$
VEE	$K - 1 + pK$	$K + (p - 1) + \frac{p(p-1)}{2}$	$\frac{(p+2)(p-1)}{2} + Kp + 2K - 2$	$p^2 + Kp + K$
EVV	$K - 1 + pK$	$1 + K(p - 1) + K\frac{p(p-1)}{2}$	$K\frac{(p+2)(p-1)}{2} + Kp + K$	$Kp^2 + Kp + K$
VVV	$K - 1 + pK$	$K + K(p - 1) + K\frac{p(p-1)}{2}$	$K\frac{(p+2)(p-1)}{2} + Kp + 2K - 1$	$Kp^2 + Kp + K$

Table 1.2: Full Table of Parameters

There is an attractive advantage in using the \mathbf{LDL}^\top decomposition. Since both the \mathbf{LDL}^\top and eigendecomposition derive from the same covariance matrix, the necessary parameters are the same in cardinality. In the case of the \mathbf{Q} and \mathbf{L} matrices, there need to be $\frac{p(p-1)}{2}$ parameters to be determined to uniquely define these matrices. In the case of the \mathbf{L} matrix these are straightforward the entries of the lower diagonal matrix, whereas \mathbf{Q} needs a nontrivial amount of work to determine a minimal generating set of parameters, which makes computation of the decomposition as in [Celeux and Govaert \(1995\)](#) a lot more difficult. Therefore the \mathbf{LDL}^\top decomposition was chosen for the purpose of this thesis.

1.5 Problems of the EM-algorithm

The EM-algorithm has stalling problems especially close to a local optimum. In their seminal work, [Dempster, Laird, and Rubin \(1977\)](#), have proven that the EM-algorithm converges under mild regularity conditions. However, convergence does not guarantee fast convergence. In fact, a lot of the work, that has gone into the research around the EM-algorithm has been concerned with speeding up convergence, see [McLachlan and Peel \(2000\)](#)[section 2.17]. The concern here is that a slowing in convergence might be mistaken for actual convergence.

This phenomenon is not infrequent and in difficult mixtures quite visible. To illustrate let us look at a particular mixture taken from [Marron and Wand \(1992\)](#) and the `nor1mix` package from CRAN. `nor1mix` is a package designed and developed for educational purposes to teach about univariate normal mixtures. It is also the spiritual predecessor of this thesis' R code.

The mixture is a trimodal mixture of uneven weight, as shown in figure 1.2. While not the most difficult mixture studied by [Marron and Wand \(1992\)](#), it is certainly not trivial either. In the figures below, 1.5 and 1.5, we demonstrate, that even after 200 iterations of the EM-algorithm the convergence is poor. In this instance, the initialization is done using R's CLARA implementation from the cluster package.

then an illustration of MW examples of pathological cases

We can see, that the EM-algorithm seems to converge to an intermediary solution, where the smaller middle solution is weighted lower, until it manages to correct back and find

```

> library("nor1mix")
> MW.nm9 ## Trimodal mixture
'Normal Mixture' object
      mu sigma  w
[1,] -1.2  0.60 0.45
[2,]  1.2  0.60 0.45
[3,]  0.0  0.25 0.10

```

Figure 1.1: Parameters of MW.nm9

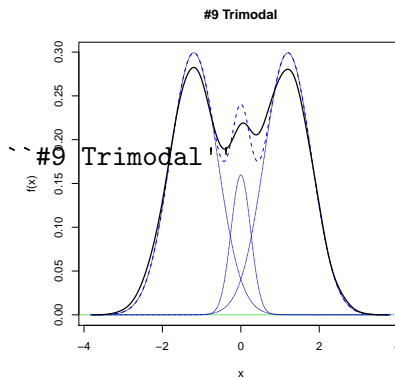


Figure 1.2: True and Estimated density

147 the correct components.




Figure 1.3: 20 EM steps




Figure 1.4: 200 EM steps

148 We see how change in log-likelihood seems to stagnate. However, this does not stay that
 149 way. If we let EM run a bit further we see, the log-likelihood hits a flatspot, after which
 150 convergence accelerates again.

151 In fact, it seems that the previous solution is a saddle point in the likelihood function,
 152 where EM has chronic problems continuing improvements.

153 give 2D demonstration.

154 1.6 Alternative Option

155 In conclusion, the EM-algorithm has very appealing advantages. However, as we have
 156 shown, there are chronic problems in convergence rates. The aim of this thesis is to test
 157 if some improvement could be achieved by a different method.

158 The plan is reasonably straightforward:

- 159 i.) Initialize using CLARA.
- 160 ii.) Perform one m-step, to transform CLARA's results into the form of a normal mixture.




Figure 1.5: Log-likelihood Plotted against Iteration Count for the Example in 1.5

- 161 iii.) Apply a general optimizer, using the mixture's log-likelihood function.
- 162 what do we hope from this? better convergence proof of concept i.e. not complete failure
- 163 raise questions about implementation, clara fctn optim params
- 164 the subsequent chapter is devoted to answering this question by documenting the devel-
- 165 opment of norMmix

Chapter 2

The norMmix Package

2.1 Introduction to the Package

For this thesis, an R package was developed that implements the algorithm that fits multivariate normal mixtures to given data. ¹ There is a lot of unused code still in the package. These were at one point implemented used and discarded. They are still included for demonstration. The `norMmix` package is constructed around the `norMmix` object that codifies a normal Multivariate mixture model, and the `llnorMmix()` function, that calculates the log-likelihood given a model and data.

In the table 2.1 the notation used in code is listed along with a translation to the previously used mathematical notation. Additionally, functions with ambiguous names are listed here.

The package contains the following functionality:

The package relies on `optim` from the `stats` package for general optimization. we use the standard method implemented in `optim` which is `BFGS`, which is a quasi-Newton method (also known as a variable metric algorithm) as described in [Broyden \(1970\)](#) among others.

The workflow when using the package is as follows. The function `rnorMmix` can be used to generate data from a `norMmix` object. The MW objects provide ready made examples and

¹The package was written with R version 3.6.1 (2019-07-05) last updated on 2019-10-22.

In Notation	In Code
π_i	<code>w, weights</code>
Σ	<code>Sigma</code>
μ	<code>mu</code>
K	<code>k</code>
dimension	<code>p, dim, dims</code>
components	<code>cl, components</code>
Σ model	<code>model</code>
cluster's CLARA	<code>clara</code>
mclust's hierarchical clustering	<code>mclVVV</code>
mclust's Mclust fuction	<code>mclust</code>

Table 2.1: Translation Table: Mathematical Notation to R Code

norMmix `norMmix()` is the 'init-method' for `norMmix` objects. There exist `is.norMmix`, `rnorMmix` and `dnorMmix` functions.

parametrization The main functions that handle reparametrization of models from and to LDL^T decomposition are `nMm2par` and `par2nMm`, which are inverse to each other.

MLE The function `norMmixMLE` marries the main components of this package. It initializes a model and parametrizes it for use with `optim`

model choice Using `norMmixMLE`, the function `fitnMm` allows fitting of multiple models and components. Functions analyzing the output of this are also provided, e.g. `BIC` and `print` methods.

misc There are also various methods of generics, like `logLik`, `print`, `BIC`, `AIC` and `nobs` as well as various `print` methods.

example objects Following the paper of [Marron and Wand \(1992\)](#) various example objects are provided and used for study. They follow the naming convention: MW + dimension + number. For example MW213 for the 13th model of dimension 2.

simulations A good portion of the package is designed with the study of simulations in mind. Therefore there are functions provided to study large collections of evaluated data. e.g. `complot`

183 objects of study and the `norMmix` function can be used to define normal mixtures from
 184 scratch. Of course, other data sets can be used for analysis. The following functions rely,
 185 however, on the `matrix` data structure. So dataframes must be converted beforehand and
 186 non numerical data is not accepted.

187 Given data, the functions that accept it for analysis are mainly `norMmixMLE` and `fitnMm`.
 188 The former performs model fit on data, and the latter performs model selection, by calling
 189 `norMmixMLE` for specified `k` and `model` vectors.

190 2.1.1 `norMmixMLE`

191 The core of `norMmixMLE` is the application of `optim` in conjunction with `llnorMmix` as
 192 function to be optimized. `llnorMmix` can be accessed directly, however, it needs a trans-
 193 posed dataset. As stated in section 1.6 the MLE implicitly performs initialization. There
 194 are two options for this initialization step. One is the CLARA clustering algorithm, with
 195 non-standard arguments. The standard arguments are somewhat historic in origin and
 196 were, at the time, chosen because of hardware limitations. The newer function, due to
 197 this thesis' advisor Martin Mächler, was designed to be a 'sensible' alternative, but should
 198 be subject to further scrutiny. It is reproduced here.

```
> norMmix:::ssClaraL
function (n, k, p)
pmin(n, pmax(40, round(10 * log(n))) + round(2 * k * pmax(1,
  log(n * p))))
<bytecode: 0x4903608>
<environment: namespace:norMmix>
```

199 It is dependent on the size and dimension of the dataset, as well as the demanded number
 200 of clusters. The alternative to CLARA is `mclust`'s hierarchical agglomerative clustering,
 201 which follows the work of [Fraley \(1998\)](#). The intention behind using `mclust`'s initialization
 202 function is to directly compare how much difference the initialization process makes.

The initialization stage does not yield a normal mixture. This requires a way to transform a clustering into a mixture. The method chosen in this package is to use an m-step from the EM-algorithm. Unlike the EM-algorithm, clustering algorithms like CLARA produce binary cluster membership results, whereas the component membership of EM is determined as a probability value between 0 and 1. This is resolved by interpreting the results as probability values which are either 0 or 1. These are then used as the τ_j as described in section 1.2. This m-step is also taken from the `mclust` package for reasons better explained in section 2.2. It has the advantage of being able to generate a mixture object with the correct covariance model.

This mixture object is still in human readable form and not the necessary parameter vector demanded by `optim`. So an application of the function `nMm2par` is carried out, resulting in a starting value for `optim`.

Due to the nature of the package the returned results are more than abundant. Not only is the fitted model returned but also everything produced by `optim` and the entire dataset. Here are listed the stucture the returned values:

```
> data(fSMI.12, package="norMmix")
> str(fSMI.12$nMm[3,3][[1]], max=2)
```

List of 6

```
$ norMmix:List of 6
..$ mu      : num [1:20, 1:3] 15.9 30.7 36.2 21.8 753 ...
..$ Sigma   : num [1:20, 1:20, 1:3] 0.358 0 0 0 0 ...
..$ weight  : num [1:3] 0.219 0.419 0.362
..$ k       : int 3
..$ dim     : int 20
..$ model   : chr "EEI"
..- attr(*, "name")= chr "model = EEI , clusters = 3"
..- attr(*, "class")= chr "norMmix"
$ optr      :List of 5
..$ par      : num [1:82] 0.264 0.118 15.903 30.67 36.155 ...
..$ value    : num 7370
..$ counts   : Named int [1:2] 232 88
.. ..- attr(*, "names")= chr [1:2] "function" "gradient"
..$ convergence: int 0
..$ message   : NULL
$ npar      : int 82
$ n         : int 141
$ x         : num [1:141, 1:20] 16.1 15.7 15.7 16.1 16.6 ...
..- attr(*, "dimnames")=List of 2
$ cond      : num 1.72
- attr(*, "class")= chr "norMmixMLE"
```

Besides `mclust` the package also relies on a number of other packages for various tasks. Listed in no particular order: `cluster`, `MASS`, `mvtnorm`, `mclust`, `mixtools` and `sfsmisc`.

since `mclust` is one of the more popular packages implementing the EM algo, we employ a lot of functions from `mclust`, to keep things around EM as similar as possible.

also relies on `mixtools` package for random generating function `rnorMmix` using `rmvnorm`.

2.2 On The Development of norMmix

about Cholesky decomp as `ldlt`. has advantages: fast, parametrically parsimonious, can easily compute loglikelihood

One dead-end was the parametrization of the weights of a mixture using the `logit` function.

```
> logit <- function(e) {
+   stopifnot(is.numeric(e) ,all(e >= 0), all.equal(sum(e),1))
+   qlogis(e[-1L])
+ }
> logitinv <- function(e) {
+   if (length(e)==0) {return(c(1))}
+   stopifnot(is.numeric(e))
+   e<- plogis(e)
+   sp. <- sum(e)
+   w <- c((1-sp.), e)
+ }
```

This uses the logistical function `logis` to transform to reduce the number of weights from K to $K - 1$. Much like `clr1`, given a list of weights `logit` will transform them and `logitinv` will correctly reverse the transformation. However, unlike `clr1`, it will not transform an arbitrary list of length $K - 1$ into a valid weight parameter. For example:

```
> w <- runif(7); ret <- logitinv(w)
> ret

[1] -3.5617264  0.6521836  0.7071940  0.5431553  0.6491777  0.7057722  0.6832753
[8]  0.6209682
```

The issue here is that the last line of `logitinv`, which is necessary to sum to one, but results in a negative value in `ret[1]` which is not a valid weight. The underlying issue is that not every tuple in \mathbb{R}^{K-1} is a result of `logit`.

The option to use `logit` is still an argument to `norMmixMLE` by specifying `trafo="logit"`, but it shouldn't be used.

Another issue during development cropped up during fitting of high dimensional data. We studied the dataset `SMI.12` from the package `copula`:

```
> data(SMI.12, package="copula")
> str(SMI.12)

num [1:141, 1:20] 16.1 15.7 15.7 16.1 16.6 ...
- attr(*, "dimnames")=List of 2
 ..$ : chr [1:141] "2011-09-09" "2011-09-12" "2011-09-13" "2011-09-14" ...
 ..$ : chr [1:20] "ABBN" "ATLN" "ADEN" "CSGN" ...
```

A consequence of high dimensions is that matrix multiplication is no longer very stable. As a result, the covariance matrices produced by our own implementation of the EM-algorithms `m-step` (`mstep.nMm`) were not positive definite. In the case of `SMI.12`, several covariance matrices are degenerate, which results in cancellation error with near-zero entries. We attempted to correct this with the function `forcePositive`, which simply tries

```
> plot(MW215)
```

Figure 2.1: Demonstration of the MW Objects

to set D in LDL^\top greater than zero. This didn't resolve the issue, since a non-negligible part of the numerical error was in the L matrix and the resultant covariance matrix was still not positive definite.

We eventually resolved this issue by abandoning our own implementation and using the functions from the `Mclust` package. Not only were these numerically stable they were also able to differentiate between models, whereas ours would assume VVV for every fit.

testing of mvtnorm as proof that ldlt is in fact faster parametrization

mention, that there may be faster ways to apply backsolve. quote knuth about premature optimization?

not possible to sensibly compare normal mixtures except maybe a strange sorting algorithm using Mahalanobis distance or Kullback-Leibler distance or similar (Hellinger), but not numerically sensible to integrate over potentially high-dimensional spaces.

2.3 Demonstration

Mention, that `mclust` doesn't depend on seed(double check) and therefore `norMmix` has 'advantage' of 'confidence intervals'. We can run 50 simulations and see if there might be more sensible clusters.

```
> set.seed(2019); x <- rnorMmix(500, MW215)
> mleResult <- norMmixMLE(x, 3, "VEE")

initial value 2206.907425
iter 10 value 2147.633703
iter 20 value 2125.658743
final value 2125.658364
converged

> mleResult

object of class 'norMmixMLE'
norMmix object:
multivariate normal mixture model with the following attributes:
name:                model = VEE , components = 3
dimension:           2
components:           3
weight of components 0.365 0.325 0.31

returned from optim:
function gradient
      75      22

log-likelihood: -2125.658
```




Figure 2.2: The Correct Mixture




Figure 2.3: The Fitted Mixture

nobs	npar	nobs/npar
500	13	38.46154

Chapter 3

Comparing Algorithms

With the `norMmix` package explained, we can turn to comparing it to existing methods. As previously stated, the implementation representing the EM-algorithm is the `mclust` package. It will be used with very little deviation from out-of-the-box, safe for restriction to the covariance models. This is done, so we can compare like with like. The specific command that performs the EM-algorithm is:

```
> #mclust::Mclust(x, G=c1, modelNames=mo)$BIC
```

Where `c1` is a vector of integers of however many components we are trying to fit and `mo` are the model names:

```
[1] "EII" "VII" "EEI" "VEI" "EVI" "VVI" "EEE" "VEE" "EVV" "VVV"
```

The `$BIC` element of the results is taken as the main tool for model selection, as it is advertised in the package authors paper [Scrucca et al. \(2016\)](#).

There is however a small but crucial change applied to these results. The `mclust` package authors have flipped the definition of the BIC to mean:

$$2\ln(\hat{L}) - \ln(n)\theta$$

instead of the more common

$$\ln(n)\theta - 2\ln(\hat{L})$$

Where n is the number of observations, θ is the cardinality of the parameter vector and \hat{L} is the estimated log-likelihood.

So even if not explicitly mentioned, we use the negative of the values returned by `mclust`. here show bic type plots, how to read them, and what we're trying to compare i.e. `clara`, `mclVVV`, `mclust`.

First, we illustrate the structure of the graphical results we will be presenting hereafter. The basic shape of the plots will be the BIC value plotted against the number of components. This is in line with `mclust`'s manner of visualizing data, however since our method is to some extent RNG dependent, we are forced to display multiple runs of the algorithm on the same graph. Therefore we split the plot according to covariance model, putting 10 models in 10 graphs in a plot. Here an example:

Example Plot




Figure 3.1: Example of Comparison Plot

As can be seen from the formula of the BIC value, lower is better. When selecting a model based on BIC, we take the model and component with the lowest value to be the best fitting model. Although this may not necessarily be the 'correct' model, that is, the model from which the data arises.

There are many ways in which this type of model selection might miss the correct model, for example by 'gluing together' multiple components into one, or covering the dataset in a 'patchwork' of smaller components, to name a few.

We will discuss them as they arise in the following analysis of simulations

here explain simulations conducted, [A.2](#) here explain the various sections: time, n, p, difficult, nonnormal

3.1 Time Analysis

here how much time they take, in p,k and n give approximate $O(x)$ value

```
> library(norMmix, lib.loc="~/ethz/BA/norMmix.Rcheck/")
> # change this dir to wherever the simulations are saved
> mainsav <- normalizePath("~/ethz/BA/Rscripts/")
> savdir <- file.path(mainsav, "2time")
> filelist <- list.files(savdir, pattern=".rds")
> filelist <- grep("mcl.rds", filelist, invert=TRUE, value=TRUE)
> ## need to split these better
> f <- lapply(file.path(savdir, filelist), function(j) readRDS(j)$fit)
> times <- unlist(lapply(f, function(j) extracttimes(j)[,1]))
> dims <- unlist(lapply(f, function(j) attr(extracttimes(j), "p")))
> size <- unlist(lapply(f, function(j) attr(extracttimes(j), "n")))
> ddims <- rep(dims, each=80)
> ssize <- rep(size, each=80)
> pars <- unlist(lapply(f, npar))
> r <- lm(log(times) ~ log(pars) + log(ddims) + log(ssize))
> summary(r)
```

Call:

```
lm(formula = log(times) ~ log(pars) + log(ddims) + log(ssize))
```

Residuals:

	Min	1Q	Median	3Q	Max
	-3.4428	-0.2986	0.0671	0.4579	2.0936

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-9.74133	0.10598	-91.91	<2e-16 ***
log(pars)	2.75983	0.01181	233.75	<2e-16 ***
log(ddims)	-2.06063	0.02483	-82.99	<2e-16 ***
log(ssize)	0.61301	0.01446	42.38	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6946 on 7196 degrees of freedom
 Multiple R-squared: 0.8887, Adjusted R-squared: 0.8887
 F-statistic: 1.916e+04 on 3 and 7196 DF, p-value: < 2.2e-16

```
> plot(times~pars, log="xy", yaxt="n", xaxt="n")
> sfsmisc::eaxis(1)
> sfsmisc::eaxis(2)
```




Figure 3.2: Log-log Plot of System Time against Parameter Length

297 can see that time is almost one to one proportional to parameter length.

298 3.2 Behaviour in n

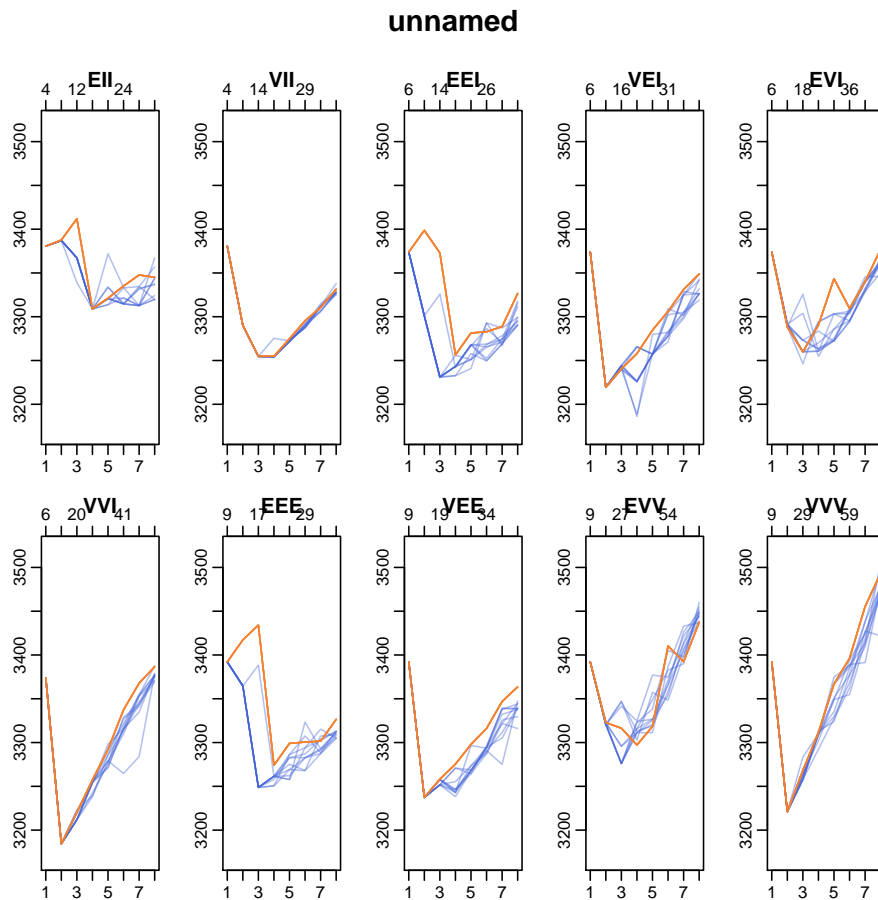
299 here show as expected narrower scattering as n increases

300 [h]

301 3.3 Behaviour in p

302 here show how norMmix is consistently competitive with mclust


```
> compplot(s05mw34bic, m0534)
```



3.4 Difficult Mixtures

here show behaviour in difficult cases


```
> savdir <- file.path(mainsav, "2init")
> filenames <- list.files(savdir, pattern=".rds")
> MW214fn <- grep("MW214", filenames, value="TRUE")
> mclustfiles <- grep("mcl.rds", MW214fn, value=TRUE)
> MW214fn <- grep("mcl.rds", MW214fn, value="TRUE", invert=TRUE)
> claraMW <- grep("clara", MW214fn, value=TRUE)
> mclMW <- grep("mclVVV", MW214fn, value=TRUE)
> clarabic <- massbic(claraMW, savdir)
> mclbic <- massbic(mclMW, savdir)
> mclustbic <- readRDS(file.path(savdir, mclustfiles[1]))
```

here some examples of fitted mixtures


We can see, that, subtracting the obvious hiccups of the small erroneous components, `norMmix` has correctly found the 'intended' distribution. This is remarkable, given the small sample size and difficulty of distribution

```
> compplot(s10mw34bic, m1034)
```

unnamed



```
> plot(MW34)
```



3.5 Nonnormal Mixtures

here 2smi and 2var, maybe others as well.

here 2smi:

```
> savdir <- file.path(mainsav, "2smi")
> filenames <- list.files(savdir, pattern=".rds")
> fnclara <- grep("clara_seed", filenames, value=TRUE)
> fnmclVV <- grep("mclVVV_see", filenames, value=TRUE)
> fnmclus <- grep("__mcl.rds", filenames, value=TRUE)
```


While not very spectacular, the graphs show that even at large parameter counts our algorithm closes in on the same values as `mclust`. At these dimensions it is difficult to compare if these are actually equal, or even similar fits, but going by BIC values, it is at the very least equally viable as a working model.

To illustrate, here are the parameter sizes for this simulation:

	EII	VII	EEI	VEI	EVI	VVI	EEE	VEE	EVV	VVV
1	21	21	40	40	40	40	230	230	230	230
2	42	43	61	62	80	81	251	252	460	461
3	63	65	82	84	120	122	272	274	690	692
4	84	87	103	106	160	163	293	296	920	923

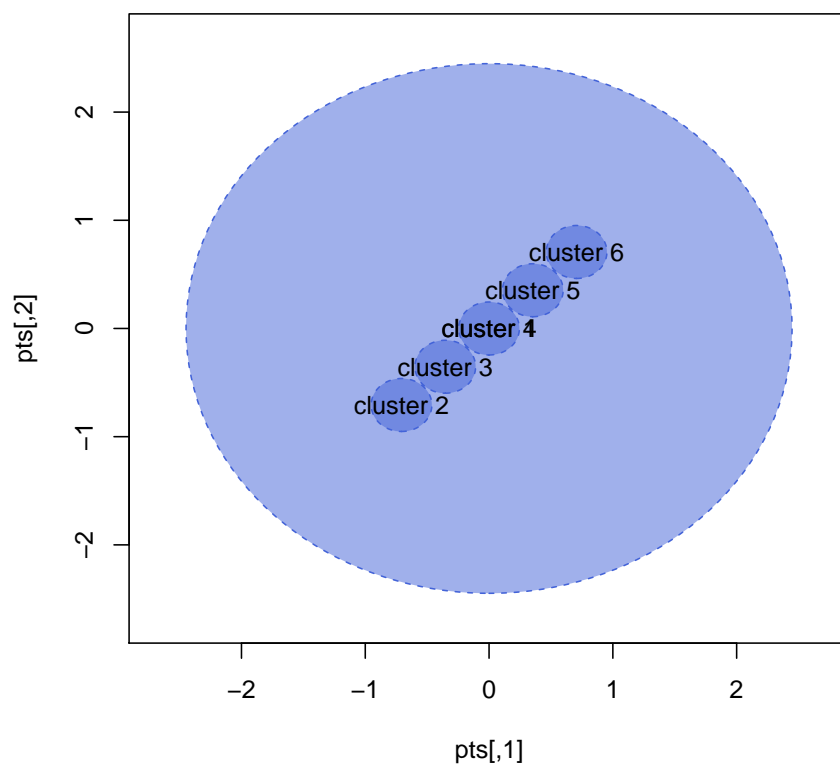
```
> complot(clarabic, mclbic, mclustbic, main="Fit of MW34")
```

Fit of MW34




```
5 105 109 124 128 200 204 314 318 1150 1154
6 126 131 145 150 240 245 335 340 1380 1385
7 147 153 166 172 280 286 356 362 1610 1616
8 168 175 187 194 320 327 377 384 1840 1847
```

```
> plot(MW214)
```

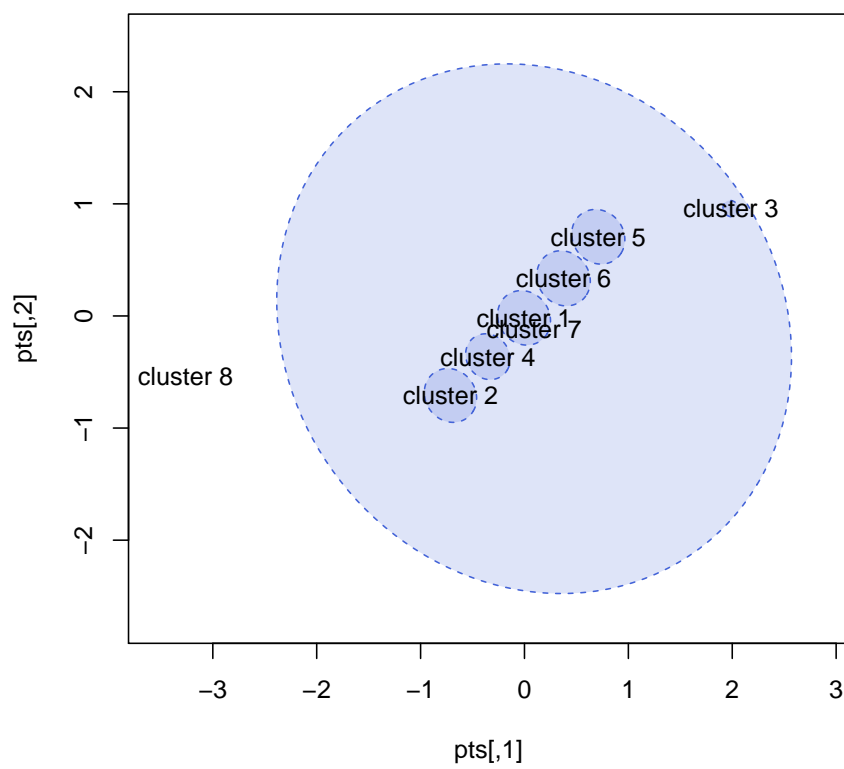


```
> compplot(clarabic, mclbic, mclustbic, main="Fit of MW214")
```


Fit of MW214



```
> f <- readRDS(file.path(savdir, claraMW[28]))  
> ff <- f$fit$nMm[8,8][[1]]  
> plot(ff$norMmix)  
> #points(ff$x)
```



BIC of SMI.12



317 Chapter 4

318 Discussion

319 one shortcoming is time inefficiency. largely due to implementation. mclust has 16'000
320 lines of Fortran code, impossible in the scope of this thesis.

321 proof of concept?? definitely possible to do model selection using a general optimizer.

322 strong points: 'randomness' of clara/optim allows 'confidence intervalls' for selected model
323 flexibility of approach: given an logLik fctn can do mixture fitting w/ arbitrary models

324 further study might include: other presumed component distributions, 'high' dimensions

325 failures of implementation: no lower boundary for variance, can lead to minuscule compo-
326 nents not as bad as it used to be.

Bibliography

- Benaglia, T., D. Chauveau, D. R. Hunter, and D. Young (2009). mixtools: An R package for analyzing finite mixture models. *Journal of Statistical Software* 32(6), 1–29.
- Broyden, C. G. (1970, 03). The Convergence of a Class of Double-rank Minimization Algorithms 1. General Considerations. *IMA Journal of Applied Mathematics* 6(1), 76–90.
- Celeux, G. and G. Govaert (1995). Gaussian parsimonious clustering models. *Pattern Recognition* 28(5), 781–793.
- Dempster, A., N. Laird, and D. Rubin (1977). Maximum likelihood from incomplete data via the em algorithm. *Journal of the Royal Statistical Society: Series B (Methodological)* 39(1), 1–22.
- Fraley, C. (1998). Algorithms for model-based gaussian hierarchical clustering. *SIAM Journal on Scientific Computing* 20(1), 270–281.
- Genz, A., F. Bretz, T. Miwa, X. Mi, F. Leisch, F. Scheipl, and T. Hothorn (2019). *mvt-norm: Multivariate Normal and t Distributions*. R package version 1.0-11.
- Maechler, M. (2019). *sfsmisc: Utilities from 'Seminar fuer Statistik' ETH Zurich*. R package version 1.1-4.
- Maechler, M., P. Rousseeuw, A. Struyf, M. Hubert, and K. Hornik (2019). *cluster: Cluster Analysis Basics and Extensions*. R package version 2.1.0 — For new features, see the 'Changelog' file (in the package source).
- Marron, S. and M. Wand (1992). Exact mean integrated squared error. *The Annals of Statistics* 20(2), 712–736.
- McLachlan, G. and D. Peel (2000). *Finite Mixture Models* (1 ed.). Wiley Series in Probability and Statistics. Wiley-Interscience.
- Pearson, K. and O. M. F. E. Henrici (1896). Vii. mathematical contributions to the theory of evolution. iii. regression, heredity, and panmixia. *Philosophical Transactions of the Royal Society of London. Series A, Containing Papers of a Mathematical or Physical Character* 187, 253–318.
- Scrucca, L., M. Fop, B. Murphy, and A. Raftery (2016). mclust 5: Clustering, classification and density estimation using gaussian finite mixture models. *R Journal* (8(1)), 289–317.
- Venables, W. N. and B. D. Ripley (2002). *Modern Applied Statistics with S* (Fourth ed.). New York: Springer. ISBN 0-387-95457-0.

359 Appendix A

360 R Code

361 A.1 llnorMmix

362 Here llnorMmix, since it is the central piece of the package, and 2time.R as an example
363 of a simulation script.

```
##### the llnorMmix function, calculating log likelihood for a given
##### parameter vector
## Author: Nicolas Trutmann 2019-07-06
## Log-likelihood of parameter vector given data
#
# par:    parameter vector
# tx:     transposed sample matrix
# k:      number of components
# model:  assumed distribution model of normal mixture
# trafo:  either centered log ratio or logit
llnorMmix <- function(par, tx, k,
                      trafo=c("clr1", "logit"),
                      model=c("EII","VII","EEI","VEI","EVI",
                              "VVI","EEE","VEE","EVV","VVV")
                      ) {
  stopifnot(is.matrix(tx),
            length(k <- as.integer(k)) == 1, k >= 1)
  p <- nrow(tx)
#   x <- t(x) ## then only needed in (x-mu[,i])^2 i=1..k
# 2. transform
  model <- match.arg(model)
  trafo <- match.arg(trafo)
  l2pi <- log(2*pi)
# 3. calc log-lik
# get w
  w <- if (k==1) 1
    else switch(trafo,
                "clr1" = clr1inv (par[1:(k-1)]),
                "logit"= logitinv(par[1:(k-1)]),
```

```

        stop("invalid 'trafo': ", trafo)
    )
    # start of relevant parameters:
    f <- k + p*k # weights -1 + means +1 => start of alpha
    # get mu
    mu <- matrix(par[k:(f-1L)], p,k)
    f1 <- f      # end of alpha if uniform
    f2 <- f+k-1L # end of alpha if var
    f1.1 <- f1 +1L # start of D. if alpha unif.
    f2.1 <- f1 + k # start of D. if alpha variable
    f11 <- f1 + p-1 # end of D. if D. uniform and alpha uniform
    f12 <- f1 +(p-1)*k # end D. if D. var and alpha unif.
    f21 <- f2 + p-1 # end of D. if D. uniform and alpha variable
    f22 <- f2 +(p-1)*k # end of D. if D. var and alpha var.
    f11.1 <- f11 +1L # start of L if alpha unif D unif
    f21.1 <- f21 +1L # start of L if alpha var D unif
    f12.1 <- f12 +1L # start of L if alpha unif D var
    f22.1 <- f22 +1L # start of L if alpha var D var
    f111 <- f11 + p*(p-1)/2 # end of L if alpha unif D unif
    f211 <- f21 + p*(p-1)/2 # end of L if alpha var D unif
    f121 <- f12 + k*p*(p-1)/2 # end of L if alpha unif D var
    f221 <- f22 + k*p*(p-1)/2 # end of L if alpha var D var
    # initialize f(tx_i) i=1..n vector of density values
    invl <- 0
    # calculate log-lik, see first case for explanation
    switch(model,
    "EII" = {
        alpha <- par[f]
        invalpha <- exp(-alpha)# = 1/exp(alpha)
        for (i in 1:k) {
            rss <- colSums(invalpha*(tx-mu[,i])^2)
            # this is vector of length n=sample size
            # calculates (tx-mu)t * Sigma^-1 * (tx-mu) for diagonal
            # cases.
            invl <- invl+w[i]*exp(-0.5*(p*(alpha+l2pi)+rss))
            # adds likelihood of one component to invl
            # the formula in exp() is the log of likelihood
            # still of length n
        }
    },
    # hereafter differences are difference in dimension in alpha and D.
    # alpha / alpha[i] and D. / D.[,i]
    "VII" = {
        alpha <- par[f:f2]
        for (i in 1:k) {
            rss <- colSums((tx-mu[,i])^2/exp(alpha[i]))
            invl <- invl+w[i]*exp(-0.5*(p*(alpha[i]+l2pi)+rss))
        }
    },
    )

```

```

"EEI" = {
  alpha <- par[f]
  D. <- par[f1.1:f11]
  D. <- c(-sum(D.),D.)
  D. <- D.-sum(D.)/p
  invD <- exp(alpha+D.)
  for (i in 1:k) {
    rss <- colSums((tx-mu[,i])^2/invD)
    invl <- invl+w[i]*exp(-0.5*(p*(alpha+l2pi)+rss))
  }
},
"VEI" = {
  alpha <- par[f:f2]
  D. <- par[f2.1:f21]
  D. <- c(-sum(D.), D.)
  D. <- D.-sum(D.)/p
  for (i in 1:k) {
    rss <- colSums((tx-mu[,i])^2/exp(alpha[i]+D.))
    invl <- invl+w[i]*exp(-0.5*(p*(alpha[i]+l2pi)+rss))
  }
},
"EVI" = {
  alpha <- par[f]
  D. <- matrix(par[f1.1:f12],p-1,k)
  D. <- apply(D.,2, function(j) c(-sum(j), j))
  D. <- apply(D.,2, function(j) j-sum(j)/p)
  for (i in 1:k) {
    rss <- colSums((tx-mu[,i])^2/exp(alpha+D.[,i]))
    invl <- invl+w[i]*exp(-0.5*(p*(alpha+l2pi)+rss))
  }
},
"VVI" = {
  alpha <- par[f:f2]
  D. <- matrix(par[f2.1:f22],p-1,k)
  D. <- apply(D.,2, function(j) c(-sum(j), j))
  D. <- apply(D.,2, function(j) j-sum(j)/p)
  for (i in 1:k) {
    rss <- colSums((tx-mu[,i])^2/exp(alpha[i]+D.[,i]))
    invl <- invl+w[i]*exp(-0.5*(p*(alpha[i]+l2pi)+rss))
  }
},
# here start the non-diagonal cases. main difference is the use
# of backsolve() to calculate tx^t Sigma^-1 tx, works as follows:
# assume Sigma = L D L^t, then Sigma^-1 = (L^t)^-1 D^-1 L^-1
# y = L^-1 tx => tx^t Sigma^-1 tx = y^t D^-1 y
# y = backsolve(L., tx)
"EEE" = {
  alpha <- par[f]
  D. <- par[f1.1:f11]

```

```

D. <- c(-sum(D.), D.)
D. <- D.-sum(D./p)
invD <- exp(alpha+D.)
L. <- diag(1,p)
L.[lower.tri(L., diag=FALSE)] <- par[f11.1:f111]
for (i in 1:k) {
  rss <- colSums(backsolve(L.,(tx-mu[,i]), upper.tri=FALSE)^2/invD)
  invl <- invl+w[i]*exp(-0.5*(p*(alpha+l2pi)+rss))
}
},
"VEE" = {
  alpha <- par[f:f2]
  D. <- par[f2.1:f21]
  D. <- c(-sum(D.), D.)
  D. <- D.-sum(D./p)
  L. <- diag(1,p)
  L.[lower.tri(L., diag=FALSE)] <- par[f21.1:f211]
  for (i in 1:k) {
    rss <- colSums(backsolve(L., (tx-mu[,i]), upper.tri=FALSE)^2/exp(alpha[i]+D.))
    invl <- invl+w[i]*exp(-0.5*(p*(alpha[i]+l2pi)+rss))
  }
},
"EVV" = {
  alpha <- par[f]
  D. <- matrix(par[f1.1:f12],p-1,k)
  D. <- apply(D.,2, function(j) c(-sum(j), j))
  D. <- apply(D.,2, function(j) j-sum(j)/p)
  L.temp <- matrix(par[f12.1:f121],p*(p-1)/2,k)
  for (i in 1:k) {
    L. <- diag(1,p)
    L.[lower.tri(L., diag=FALSE)] <- L.temp[,i]
    rss <- colSums(backsolve(L., (tx-mu[,i]), upper.tri=FALSE)^2/exp(alpha+D.))
    invl <- invl+w[i]*exp(-0.5*(p*(alpha+l2pi)+rss))
  }
},
"VVV" = {
  alpha <- par[f:f2]
  D. <- matrix(par[f2.1:f22],p-1,k)
  D. <- apply(D.,2, function(j) c(-sum(j), j))
  D. <- apply(D.,2, function(j) j-sum(j)/p)
  invalpha <- exp(rep(alpha, each=p)+D.)
  L.temp <- matrix(par[f22.1:f221],p*(p-1)/2,k)
  L. <- diag(1,p)
  for (i in 1:k) {
    L.[lower.tri(L., diag=FALSE)] <- L.temp[,i]
    rss <- colSums(backsolve(L., (tx-mu[,i]), upper.tri=FALSE)^2/invalpha[,i])
    invl <- invl+w[i]*exp(-0.5*(p*(alpha[i]+l2pi)+rss))
  }
},

```



```

    ## otherwise
    stop("invalid model:", model)
  )
  ## return  $\sum_{i=1}^n \log( f(tx_i) )$  :
  sum(log(invl))
}

sllnorMmix <- function(x, obj, trafo=c("clr1", "logit")) {
  stopifnot(is.character(model <- obj$model))
  trafo <- match.arg(trafo)
  llnorMmix(nMm2par(obj, model=model),
            tx = t(x), k = obj$k,
            model=model, trafo=trafo)
}

## log-likelihood function relying on mvtnorm function
#
# par:   parameter vector as calculated by nMm2par
# x:     matrix of samples
# k:     number of cluster
# trafo: transformation of weights
# model: assumed model of the distribution
llmvtnorm <- function(par, x, k,
                      trafo=c("clr1", "logit"),
                      model=c("EII", "VII", "EEI", "VEI", "EVI",
                              "VVI", "EEE", "VEE", "EVV", "VVV"))
  ) {
  stopifnot(is.matrix(x),
            length(k <- as.integer(k)) == 1, k >= 1)
  model <- match.arg(model)
  trafo <- match.arg(trafo)
  p <- ncol(x)
  nmm <- par2nMm(par, p, k, model=model, trafo=trafo)
  ## FIXME (speed!): dmvnorm(*, sigma= S) will do a chol(S) for each component
  ## ----- *instead* we already have LDL' and chol(S) = sqrt(D) L' !!
  ## another par2*() function should give L and D, or from that chol(Sigma), rather than S
  w <- nmm$w
  mu <- nmm$mu
  sig <- nmm$Sigma
  y <- 0
  for (i in 1:k) {
    y <- y + w[i]*mvtnorm::dmvnorm(x, mean=mu[,i], sigma=sig[,i])
  }
  sum(log(y))
}

```

364 A.2 Example Simulation Script

365 here e.g. 2init.R and write some remarks on it.

```
## Intent: analyse time as function of p,k,n
nmmdir <- normalizePath("~/BachelorArbeit/norMmix.Rcheck/")
savdir <- normalizePath("~/BachelorArbeit/Rscripts/2time")
stopifnot(dir.exists(nmmdir), dir.exists(savdir))
library(norMmix, lib.loc=nmmdir)
library(mclust)
## at n=500,p=2 can do about 250xfitnMm(x,1:10) in 24h
seeds <- 1:10
sizes <- c(500, 1000, 2000)
nmm <- list(MW214, MW34, MW51)
## => about 100 cases
# for naming purposes
nmnames <- c("MW214", "MW34", "MW51")
sizenames <- c("500", "1000", "2000")
files <- vector(mode="character")
for (nm in 1:3) {
  for (size in sizes) {
    set.seed(2019); x <- rnorMmix(size, nmm[[nm]])
    for (seed in seeds) {
      set.seed(2019+seed)
      r <- tryCatch(fitnMm(x, k=1:8,
                          optREPORT=1e4, maxit=1e4),
                    error = identity)
      filename <- sprintf("%s_size=%0.4d_seed=%0.2d.rds",
                          nmnames[nm], size, seed)
      files <- append(files, filename)
      cat("===> saving to file:", filename, "\n")
      saveRDS(list(fit=r), file=file.path(savdir, filename))
    }
  }
}
fillis <- list()
for (i in seq_along(sizes)) {
  for (j in seq_along(nmnames)) {
    # for lack of AND matching, OR match everything else and invert
    ret <- grep(paste(sizenames[-i], nmnames[-j], sep="|"),
                files, value=TRUE, invert=TRUE)
    fillis[[paste0(sizenames[i], nmnames[j])]] <- ret
  }
}
epfl(fillis, savdir)
```

366 **Appendix B**

367 **Further Plots**

368 here further plots:

369 **B.1 Chapter 3**

370 not nec best way to go about it, section differently.

371 dsfasdf

Declaration of Originality

The signed declaration of originality is a component of every semester paper, Bachelor's thesis, Master's thesis and any other degree paper undertaken during the course of studies, including the respective electronic versions.

Lecturers may also require a declaration of originality for other written papers compiled for their courses.

I hereby confirm that I am the sole author of the written work here enclosed and that I have compiled it in my own words. Parts excepted are corrections of form and content by the supervisor.

Title of work (in block letters):

Authored by (in block letters):

For papers written by groups the names of all authors are required.

Name(s):

First name(s):

Master	Student

With my signature I confirm that

- I have committed none of the forms of plagiarism described in the Citation etiquette information sheet.
- I have documented all methods, data and processes truthfully.
- I have not manipulated any data.
- I have mentioned all persons who were significant facilitators of the work.
- I am aware that the work may be screened electronically for plagiarism.
- I have understood and followed the guidelines in the document *Scientific Works in Mathematics*.

Place, date:

Signature(s):

Zürich August 19th 2009	bla

For papers written by groups the names of all authors are required. Their signatures collectively guarantee the entire content of the written paper.