



Swiss Federal Institute of Technology Zurich

Seminar for
Statistics

Department of Mathematics

Bachelor Thesis

Winter 2019

Nicolas Trutmann

Comparison of EM-algorithm and MLE using
Cholesky decomposition

Submission Date: placeholder

Advisor: placeholder

Abstract

The intent of this work is to compare The EM algorithm to a MLE approach in the case of multivariate normal mixture models using the Cholesky decomposition. The EM algorithm is widely used in statistics and is proven to converge, however in pathological cases convergence slows down considerably.

methods(not done)

results(not done)

Contents

1	Introduction to normal mixture models	1
1.1	Definitions	1
1.2	The EM-Algorithm in Sketch	2
1.3	Choice of Notation	2
1.4	Models of Covariance Matrices	3
1.5	Problems of the EM-algorithm	6
2	The <code>norMmix</code> Package	9
2.1	Introduction to the Package	9
2.2	On The Development of <code>norMmix</code>	10
2.3	Demonstration	11
3	Comparing Algorithms	13
3.1	Time Analysis	15
3.2	Behaviour in <code>n</code>	16
3.3	Behaviour in <code>p</code>	16
3.4	Diffixult Mixtures	17
3.5	Nonnormal Mixtures	19
4	Discussion	25
	Bibliography	26
A	R Code	29
A.1	<code>llnorMmix</code>	29
A.2	Example Simulation Script	34
B	Further Plots	35
B.1	Ch3	35

List of Figures

1.1	Parameters of <code>MW.nm9</code>	6
1.2	True and Estimated density	6
1.3	20 EM steps	7
1.4	200 EM steps	7
3.1	Example of Comparison Plot	14
3.2	Log-log Plot of System Time against Parameter Length	16

List of Tables

1.1	Table of Parameters of the Covariance Matrices	5
1.2	Full Table of Parameters	6
2.1	Translation Table: Mathematical Notation to R Code	9

Chapter 1

Introduction to normal mixture models

1.1 Definitions

A good and thorough introductory book is the work of [McLachlan and Peel \(2000\)](#) and the reader is encouraged to study it to learn in depth about normal mixtures and clustering. We will here give a short overview of normal mixtures to fix notation and nomenclature. The motivating idea behind mixture models is, that in real world examples a sample might be suspected to arise from more than one population or be more simply modelled by several overlaid distributions. The example of this, that is generally considered to be the first of this kind, is the one by Karl Pearson, who fitted two normal distributions with different means and variances. In his book, [Pearson and Henrici \(1896\)](#)[Section 4.d.; page 266], Pearson analyzed measurements of forehead to body length of crabs sampled from the bay of Naples. His mixture model-based approach suggested, that the crabs were evolving into two new subspecies.

While the theory of mixture models holds for a much broader class of distributions, we restrict ourselves here to the case of normal distributions, because this restriction fits more comfortably into the scope of this work and because normal distributions allow for a parsimonious parametrization, that is of interest to study.

This parametrization is the LDL^\top decomposition, which allows a very simple parametrization and a straightforward connection between degrees of freedom and necessarily generated numerical values. This will be explained further in section 1.4.

Let $\mu \in \mathbb{R}^p$, $\Sigma \in \mathbb{R}^{p \times p}$ be symmetric positive definite and $\phi(-; \mu, \Sigma)$ be the normal distribution with mean μ and covariance matrix Σ .

$\mathbf{Y}_1, \dots, \mathbf{Y}_n$

Definition 1.1.0.1. Suppose we have a random sample $\mathbf{Y}_1, \dots, \mathbf{Y}_n$ with probability density function $\mathbf{Y}_j \sim f(y_j)$ on \mathbb{R}^p . We assume that the density $f(y_j)$ of \mathbf{Y}_j can be written in the form

$$f(y_j) = \sum_{k=1}^K \pi_k \phi_k(y_j; \mu, \Sigma)$$

The π_k are called the component densities of the mixture and the ϕ_k mixture components.

For 'large' datasets there are more parsimonious parametrizations, that reduce computation time. These, for example, assume that all components have the same covariance, or have certain restrictions placed on them. We will give a detailed description of the models assumed in this thesis in section 1.4.

1.2 The EM-Algorithm in Sketch

With this definition, we immediately face the problem of how to fit these mixture components to given data. A popular algorithm to solve this problem is the Expectation-Maximization algorithm, abbreviated as EM-algorithm.

We give here a sketch of the EM-algorithm in the case of all normal mixture components.

Suppose we have a p -dimensional dataset of n samples x_1, \dots, x_n , onto which we would like to fit K normal distributions ϕ_k , $k \in 1, \dots, K$. We introduce a further explaining variable \mathbf{Z} in $\text{Mat}^{n \times K}$, with entries in $[0, 1]$ which represent the expectation that observation i belongs to component k .

The EM-algorithm is a two step, iterative process consisting of an 'e'-step and an 'm'-step. In the e-step the expectation of component membership is updated.

$$\tau_i(y_j; \Psi) = \phi_i(y_j; \mu_i, \Sigma_i) / \sum_{k=1}^K \phi_k(y_j; \mu_k, \Sigma_k)$$

and in the m-step given the component membership information we update the component means and covariances by weighted versions of the usual estimators.

$$\mu_i = \sum_{j=1}^n \tau_{ij} y_j / \sum_{j=1}^n \tau_{ij}$$

$$\Sigma_i = \sum_{j=1}^n \tau_{ij} (y_j - \mu_i)(y_j - \mu_i)^\top / \sum_{j=1}^n \tau_{ij}$$

There remains to be stated how to start the algorithm. Since both steps of the algorithm depend on data from the other, the EM-algorithm needs some form of initialization step. Most popular implementations use some form of pre clustering and use the EM-algorithm as subsequent tools to fit the data. The R-package `Mclust` for example uses hierarchical agglomerative clustering [L, M, TB, and AE. \(2016\)](#).

1.3 Choice of Notation

The classification of models in this paper relies heavily on the work of [Celeux and Govaert \(1995\)](#), however, out of necessity for clarity, we break with their notation. So as to not confuse the reader we describe here in depth the differences in notation between [Celeux and Govaert \(1995\)](#) and ours.

The basis of classification in [Celeux and Govaert \(1995\)](#) is the decomposition of a symmetric matrix into an orthogonal and a diagonal component. A symmetric positive definite

matrix Σ can be decomposed as follows

$$\Sigma = \lambda \mathbf{D} \mathbf{A} \mathbf{D}^\top$$

with \mathbf{D} an orthogonal matrix and \mathbf{A} a diagonal matrix and $\lambda = \sqrt[p]{\det(\Sigma)}$ the p -th root of the determinant of Σ .

This decomposition has an appealing geometric interpretation, with \mathbf{D} as the *orientation* of the distribution, \mathbf{A} the *shape*, and λ the *volume*. The problem of notation comes from standard conventions in linear algebra, where the letters A and D are usually occupied by arbitrary and diagonal matrices respectively. Furthermore, we intend to apply a variant of the Cholesky decomposition to Σ , the $\alpha \mathbf{L} \mathbf{D} \mathbf{L}^\top$ decomposition. This obviously raises some conflicts in notation.

Therefore we, from here on, when referring to the decomposition as described by [Celeux and Govaert \(1995\)](#), will use the following modification of notation:

$$\begin{aligned} \mathbf{D} &\mapsto \mathbf{Q} \\ \mathbf{A} &\mapsto \mathbf{\Lambda} \\ \lambda &\mapsto \alpha \\ \Sigma &= \lambda \mathbf{D} \mathbf{A} \mathbf{D}^\top = \alpha \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^\top \end{aligned}$$

These were chosen according to general conventions of linear algebra. \mathbf{Q} is usually chosen for orthonormal matrices; $\mathbf{\Lambda}$ is often a choice for diagonal matrices of eigenvectors and α was somewhat arbitrarily chosen.

1.4 Models of Covariance Matrices

As mentioned above, there are ways to constrain the covariance matrices for computational reasons. There are instances where the resulting loss of information is seen as acceptable, for example if, through consideration of the data, that a simplified model is acceptable. Another is if the sheer size of the data makes application of full generality impossible.

- We restrict the complexity of the decomposition components
- We restrict the variability of the mixture components

Let us look at the first case. We take the decomposition of a covariance matrix as $\Sigma = \alpha \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^\top$. Of these, we can simplify the structure of \mathbf{Q} and $\mathbf{\Lambda}$, by replacing them with the identity. If we set $\mathbf{Q} = \text{Id}$, we lose the freedom of orientation and if we set $\mathbf{\Lambda} = \text{Id}$ we restrict ourselves to spherical distributions.

of course, we cannot restrict λ while letting \mathbf{q} free, since

$$\mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^\top = \mathbf{Q} \text{Id} \mathbf{Q}^\top = \text{Id}$$

The second restriction simply means we hold the decomposition fixed throughout all covariance matrices. There is however an issue with the cholesky decomposition. For 10 out

of 14 cases as defined by Celeux and Govaert (1995), there exists a canonical translation of decompositions. The 6 diagonal cases need no translation; the eigen and cholesky decomposition are equal to identity. For the non-diagonal cases note that for a given sym. pos. def. matrix Σ we have decompositions:

$$\Sigma = \alpha \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^\top \quad \Sigma = \alpha \mathbf{L} \mathbf{D} \mathbf{L}^\top$$

Since in both cases the bracketing matrices \mathbf{Q} and \mathbf{L} have determinant 1 the determinant of Σ falls entirely on α . therefore α , in these particular decompositions, is equal for both. Celeux and Govaert (1995) vary σ by either varying or holding fixed the volume (α/α_k), shape ($\mathbf{\Lambda}/\mathbf{\Lambda}_k$) and orientation (\mathbf{Q}/\mathbf{Q}_k). These 3 times 2 cases would yield the 8 out of 14 cases of non-diagonal cases. However there is no canonical transform for either variable orientation and fixed shape or fixed orientation and variable shape. The reason for this is that in the $\mathbf{L} \mathbf{D} \mathbf{L}^\top$ decomposition the lower diagonal matrix \mathbf{L} holds some of the shape of the matrix, which in the eigendecomposition is in the $\mathbf{\Lambda}$ matrix. In fact, \mathbf{L} is orthogonal if and only if $\mathbf{L} = \text{Id}_{n \times n}$. Therefore we can only decompose matrices where either both or neither shape and orientation vary. See table 1.1.

While we could in theory construct the cases $\mathbf{L} \mathbf{D}_k \mathbf{L}^\top$ and $\mathbf{L}_k \mathbf{D} \mathbf{L}_k^\top$, however they do not correspond to the desired geometric intent behind the differentiation of models and are therefore not included.

Model	Σ_k C&G	volume	shape	orientation	parameters	LDL^\top as in C&G	parameters	count
EII	αI	equal	equal	-	α	as in C&G		1
VII	$\alpha_k I$	var.	equal	-	α_k			K
EEl	αA	equal	equal	coord. axes	α, λ_i			$1 + (p-1)$
VEI	$\alpha_k A$	var.	equal	coord. axes	α_k, λ_i			$K + (p-1)$
EVI	αA_k	equal	var.	coord. axes	$\alpha, \lambda_{i,k}$			$1 + K(p-1)$
VVI	$\alpha_k A_k$	var.	var.	coord. axes	$\alpha_k, \lambda_{i,k}$			$K + K(p-1)$
EEE	$\alpha Q A Q^\top$	equal	equal	equal	$\alpha, \lambda_i, q_{i,j}$	αLDL^\top	$\lambda, d_i, l_{i,j}$	$1 + (p-1) + \frac{p(p-1)}{2}$
EVE	$\alpha Q A_k Q^\top$	equal	var.	equal	$\alpha, \lambda_{i,k}, q_{i,j}$	doesn't exist		$1 + K(p-1) + \frac{p(p-1)}{2}$
VEE	$\alpha_k Q A Q^\top$	var.	equal	equal	$\alpha_k, \lambda_i, q_{i,j}$	$\alpha_k LDL^\top$	$\lambda_k, d_i, l_{i,j}$	$K + (p-1) + \frac{p(p-1)}{2}$
VVE	$\alpha_k Q A_k Q^\top$	var.	var.	equal	$\alpha_k, \lambda_{i,k}, q_{i,j}$			$K + K(p-1) + \frac{p(p-1)}{2}$
EEV	$\alpha Q_k A Q_k^\top$	equal	equal	var.	$\alpha, \lambda_i, q_{i,j,k}$	don't exist		$1 + (p-1) + K \frac{p(p-1)}{2}$
VEV	$\alpha_k Q_k A Q_k^\top$	var.	equal	var.	$\alpha_k, \lambda_i, q_{i,j,k}$			$K + (p-1) + K \frac{p(p-1)}{2}$
EVV	$\alpha Q_k A_k Q_k^\top$	equal	var.	var.	$\alpha, \lambda_i, q_{i,j,k}$	$\alpha L_k D_k L_k^\top$	$\lambda, d_{i,k}, l_{i,j,k} \quad j > i$	$1 + K(p-1) + K \frac{p(p-1)}{2}$
VVV	$\alpha_k Q_k A_k Q_k^\top$	var.	var.	var.	$\alpha_k, \lambda_i, q_{i,j,k}$	$\alpha_k L_k D_k L_k^\top$	$\lambda_k, d_{i,k}, l_{i,j,k} \quad j > i$	$K + K(p-1) + K \frac{p(p-1)}{2}$

Table 1.1: Table of Parameters of the Covariance Matrices

Σ model	μ, π	Σ	reduced	$\mathcal{O}()$
EII	$K - 1 + pK$	1	$pK + K$	pK
VII	$K - 1 + pK$	K	$pK + 2K - 1$	pK
EEI	$K - 1 + pK$	$1 + (p - 1)$	$pK + p + K - 1$	pK
VEI	$K - 1 + pK$	$K + (p - 1)$	$pK + p + 2K - 2$	pK
EVI	$K - 1 + pK$	$1 + K(p - 1)$	$2pK$	pK
VVI	$K - 1 + pK$	$1 + K(p - 1)$	$2pK + 1$	pK
EEE	$K - 1 + pK$	$1 + (p - 1) + \frac{p(p-1)}{2}$	$\frac{p(p-1)}{2}$	p^2
VEE	$K - 1 + pK$	$K + (p - 1) + \frac{p(p-1)}{2}$	$K + \frac{(p+2)(p-1)}{2}$	$p^2 + K$
EVV	$K - 1 + pK$	$1 + K(p - 1) + \frac{p(p-1)}{2}$	$\frac{(p+K)(p-1)}{2} + 1$	$p^2 + Kp + K$
VVV	$K - 1 + pK$	$K + K(p - 1) + \frac{p(p-1)}{2}$	$Kp + \frac{p(p-1)}{2}$	$p^2 + Kp + K$

Table 1.2: Full Table of Parameters

here explain why `ldlt` is so nice to parametrize; i.e. `Q` in `QLambdaQt` has as many degrees of freedom, but needs more complicated boundary conditions to work well.

1.5 Problems of the EM-algorithm

The EM-algorithm has stalling problems especially close to a local optimum. In their seminal work, [Dempster, Laird, and Rubin \(1977\)](#), have proven that the EM-algorithm converges under mild regularity conditions. However, convergence does not guarantee fast convergence. In fact, a lot of the work, that has gone into the research around the EM-algorithm has been concerned with speeding up convergence, see [McLachlan and Peel \(2000\)](#)[section 2.17]. In common software implementations, The concern here is that a slowing in convergence might be mistaken for actual convergence.

This phenomenon is not infrequent and in difficult mixtures quite visible. To illustrate let us look at a particular mixture taken from [Marron and Wand \(1992\)](#) and the `nor1mix` package from CRAN.

```
> library("nor1mix")
> MW.nm9 ## Trimodal mixture
'Normal Mixture' object
      mu sigma  w
[1,] -1.2  0.60 0.45
[2,]  1.2  0.60 0.45
[3,]  0.0  0.25 0.10
```


Figure 1.1: Parameters of `MW.nm9`

Figure 1.2: True and Estimated density

then an illustration of MW examples of pathological cases

here we see how change in loglik seems to stagnate. However, this does not stay that way, if we let EM run a bit further.

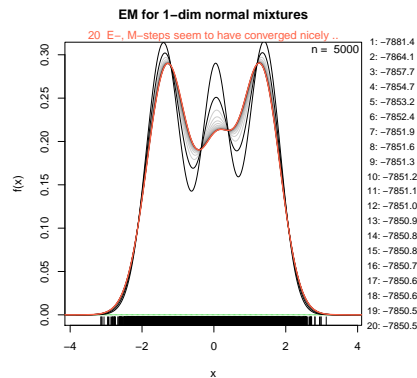


Figure 1.3: 20 EM steps




Figure 1.4: 200 EM steps

to conclude example show part of mixest that shows it takes 1200 iterations to converge

In fact, it seems that the previous solution is a saddle point in the likelihood function, where EM has chronic problems continuing improvements.

give 2D demonstration.

maybe show Marr Wand's examples of 'difficult' mixtures

give conclusion recapping the just demonstrated, and lead in for next chapter

Chapter 2

The norMmix Package

2.1 Introduction to the Package

For this thesis, an R package was developed that implements the algorithm that fits multivariate normal mixtures to given data.¹ There is a lot of unused code still in the package. These were at one point implemented used and discarded. They are still included for demonstration. The `norMmix` package is constructed around the `norMmix` object, that codifies a `normal` Multivariate mixture model, and the `llnorMmix()` function, that calculates the log-likelihood given a model and data.

The package contains the following functionality:

relies on `optim()` generic optimizer. maximizes `llnormix` by varying model parameters.

since `mclust` is one of the more popular packages implementing the EM algo, we employ a lot of functions from `mclust`, to keep things around EM as similar as possible.

Conceptually, at the start of a fitting algo, e.g. EM we need to initialize a mixture object. thereafter the paths diverge. at the heart of `norMmix`'s functionality lie the functions: `llnorMmix` and `nMm2par` which are in turn employed by `norMmixMLE` to funnel a mixture object into `optim` and give `optim` a function to optimize.

also relies on `mixtools` package for random generating function `rnorMmix` using `rmvnorm`.

¹The package was written with R version 3.6.1 (2019-07-05) last updated on 2019-10-22.

In Notation	In Code
π_i	<code>w, weights</code>
Σ	<code>Sigma</code>
μ	<code>mu</code>
K	<code>k</code>
dimension	<code>p, dim, dims</code>
components	<code>cl, components</code>

Table 2.1: Translation Table: Mathematical Notation to R Code

norMmix `norMmix()` is the 'init-method' for `norMmix` objects. There exist `is.norMmix`, `rnorMmix` and `dnorMmix` functions.

parametrization The main functions that handle reparametrization of models from and to LDL^T decomposition are `nMm2par` and `par2nMm`, which are inverse to each other.

MLE The function `norMmixMLE` marries the main components of this package. It initializes a model and parametrizes it for use with `optim`

model choice Using `norMmixMLE`, the function `fitnMm` allows fitting of multiple models and components. Functions analyzing the output of this are also provided, e.g. `BIC` and `print` methods.

misc There are also various methods of generics, like `logLik`, `print`, `BIC`, `AIC` and `nobs` as well as various `print` methods.

example objects Following the paper of [Marron and Wand \(1992\)](#) various example objects are provided and used for study. They follow the naming convention: MW + dimension + number. for example `MW213` for the 13-th model of dimension 2.

simulations The purpose of this package is to study simulations. there are functions provided to study large collections of evaluated data. e.g `epfl`

146 2.2 On The Development of norMmix

147 about Cholesky decomp as `ldlt`. has advantages: fast, parametrically parsimonious, can
148 easily compute loglikelihood

149 maybe reread section in McLachlan about accelerating EM algo

150 not possible to sensibly compare normal mixtures except maybe a strange sorting algorithm
151 using Mahalanobis distance or Kullback-Leibler distance or similar (Hellinger), but not
152 numerically sensible to integrate over potentially high-dimensional spaces.

153 general list of (not necessarily mathematical) dead-ends in the development life of the
154 `norMmix` package. argue why this is in this section?? because, as a BScT, the learning is
155 as much part of the research as the results.

156 One dead-end was the parametrization of the weights of a mixture using the `logit` func-
157 tion.

```
> logit <- function(e) {
+   stopifnot(is.numeric(e) ,all(e >= 0), all.equal(sum(e),1))
+   qlogis(e[-1L])
+ }
> logitinv <- function(e) {
+   if (length(e)==0) {return(c(1))}
+   stopifnot(is.numeric(e))
+   e<- plogis(e)
+   sp. <- sum(e)
+   w <- c((1-sp.), e)
+ }
```

158 This uses the logistical function `logis` to transform to reduce the number of weights
159 from K to $K - 1$. Much like `clr1`, given a list of weights `logit` will transform them
160 and `logitinv` will correctly reverse the transformation. However, unlike `clr1`, it will not
161 transform an arbitrary list of length $K - 1$ into a valid weight parameter. For example:


```
> w <- runif(7); ret <- logitinv(w)
> ret

[1] -3.2617309  0.6306458  0.5682759  0.5602498  0.6616009  0.6906020  0.5707690
[8]  0.5795875
```

162 The issue here is that the last line of `logitinv`, which is necessary to sum to one, but
 163 results in a negative value in `ret[1]` which is not a valid weight. The underlying issue is
 164 that not every tuple in \mathbb{R}^{K-1} is a result of `logit`.

165 The option to use `logit` is still an argument to `norMmixMLE` by specifying `trafo="logit"`,
 166 but it shouldn't be used.

167 Another issue during development cropped up during fitting of high dimensional data. We
 168 studied the dataset `SMI.12` from the package `copula`:

```
> data(SMI.12, package="copula")
> str(SMI.12)

num [1:141, 1:20] 16.1 15.7 15.7 16.1 16.6 ...
- attr(*, "dimnames")=List of 2
..$ : chr [1:141] "2011-09-09" "2011-09-12" "2011-09-13" "2011-09-14" ...
..$ : chr [1:20] "ABBN" "ATLN" "ADEN" "CSGN" ...
```

169 A consequence of high dimensions is that matrix multiplication is no longer very stable.
 170 As a result, the covariance matrices produced by our own implementation of the EM-
 171 algorithms `m-step` (`mstep.nMm`) were not positive definite. In the case of `SMI.12`, several
 172 covariance matrices are degenerate, which results in cancellation error with near-zero en-
 173 tries. We attempted to correct this with the function `forcePositive`, which simply tries
 174 to set D in LDL^T greater than zero. This didn't resolve the issue, since a non-negligible
 175 part of the numerical error was in the L matrix and the resultant covariance matrix was
 176 still not positive definite.

177 We eventually resolved this issue by abandoning our own implementation and using the
 178 functions from the `Mclust` package. Not only were these numerically stable they were also
 179 able to differentiate between models, whereas ours would assume VVV for every fit.

180 testing of `mvtnorm` as proof that `ldlt` is in fact faster parametrization

181 mention, that there may be faster ways to apply `backsolve`. quote knuth about premature
 182 optimization?

183 2.3 Demonstration

184 Mention, that `mclust` doesn't depend on `seed`(double check) and therefore `norMmix` has
 185 'advantage' of 'confidence intervals'. We can run 50 simulations and see if there might be
 186 more sensible clusters.

187 demonstrate things; essentially put `.Rd` example sections here

Chapter 3

Comparing Algorithms

With the `norMmix` package explained, we can turn to comparing it to existing methods. As previously stated, the implementation representing the EM-algorithm is the `mclust` package. It will be used with very little deviation from out-of-the-box, safe for restriction to the covariance models. This is done, so we can compare like with like. The specific command that performs the EM-algorithm is:

```
> #mclust::Mclust(x, G=c1, modelNames=mo)$BIC
```

Where `c1` is a vector of integers of however many components we are trying to fit and `mo` are the model names:

```
[1] "EII" "VII" "EEI" "VEI" "EVI" "VVI" "EEE" "VEE" "EVV" "VVV"
```

The `$BIC` element of the results is taken as the main tool for model selection, as it is advertised in the package authors paper [L et al. \(2016\)](#).

There is however a small but crucial change applied to these results. The `mclust` package authors have flipped the definition of the BIC to mean:

$$2\ln(\hat{L}) - \ln(n)\theta$$

instead of the more common

$$\ln(n)\theta - 2\ln(\hat{L})$$

Where n is the number of observations, θ is the cardinality of the parameter vector and \hat{L} is the estimated log-likelihood.

So even if not explicitly mentioned, we use the negative of the values returned by `mclust`. here show bic type plots, how to read them, and what we're trying to compare i.e. `clara`, `mclVVV`, `mclust`.

First, we illustrate the structure of the graphical results we will be presenting hereafter. The basic shape of the plots will be the BIC value plotted against the number of components. This is in line with `mclust`'s manner of visualizing data, however since our method is to some extent RNG dependent, we are forced to display multiple runs of the algorithm on the same graph. Therefore we split the plot according to covariance model, putting 10 models in 10 graphs in a plot. Here an example:

Example Plot




Figure 3.1: Example of Comparison Plot

As can be seen from the formula of the BIC value, lower is better. When selecting a model based on BIC, we take the model and component with the lowest value to be the best fitting model. Although this may not necessarily be the 'correct' model, that is, the model from which the data arises.

There are many ways in which this type of model selection might miss the correct model, for example by 'gluing together' multiple components into one, or covering the dataset in a 'patchwork' of smaller components, to name a few.

We will discuss them as they arise in the following analysis of simulations

here explain simulations conducted, [A.2](#) here explain the various sections: time, n, p, difficult, nonnormal

3.1 Time Analysis

here how much time they take, in p,k and n give approximate $O(x)$ value

```
> library(norMmix, lib.loc="~/ethz/BA/norMmix.Rcheck/")
> # change this dir to wherever the simulations are saved
> mainsav <- normalizePath("~/ethz/BA/Rscripts/")
> savdir <- file.path(mainsav, "2time")
> filelist <- list.files(savdir, pattern=".rds")
> filelist <- grep("mcl.rds", filelist, invert=TRUE, value=TRUE)
> ## need to split these better
> f <- lapply(file.path(savdir, filelist), function(j) readRDS(j)$fit)
> times <- unlist(lapply(f, function(j) extracttimes(j)[,1]))
> dims <- unlist(lapply(f, function(j) attr(extracttimes(j), "p")))
> size <- unlist(lapply(f, function(j) attr(extracttimes(j), "n")))
> ddims <- rep(dims, each=80)
> ssize <- rep(size, each=80)
> pars <- unlist(lapply(f, npar))
> r <- lm(log(times) ~ log(pars) + log(ddims) + log(ssize))
> summary(r)
```

Call:

```
lm(formula = log(times) ~ log(pars) + log(ddims) + log(ssize))
```

Residuals:

	Min	1Q	Median	3Q	Max
	-3.4428	-0.2986	0.0671	0.4579	2.0936

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-9.74133	0.10598	-91.91	<2e-16 ***
log(pars)	2.75983	0.01181	233.75	<2e-16 ***
log(ddims)	-2.06063	0.02483	-82.99	<2e-16 ***
log(ssize)	0.61301	0.01446	42.38	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6946 on 7196 degrees of freedom
 Multiple R-squared: 0.8887, Adjusted R-squared: 0.8887
 F-statistic: 1.916e+04 on 3 and 7196 DF, p-value: < 2.2e-16

```
> plot(times~pars, log="xy", yaxt="n", xaxt="n")
> sfsmisc::eaxis(1)
> sfsmisc::eaxis(2)
```




Figure 3.2: Log-log Plot of System Time against Parameter Length

225 can see that time is almost one to one proportional to parameter length.

226 3.2 Behaviour in n


227 here show as expected narrower scattering as n increases

228 [h]

229 3.3 Behaviour in p

230 here show how norMmix is consistently competitive with mclust

```
> compplot(s05mw34bic, m0534)
```



3.4 Difficult Mixtures

here show behaviour in difficult cases

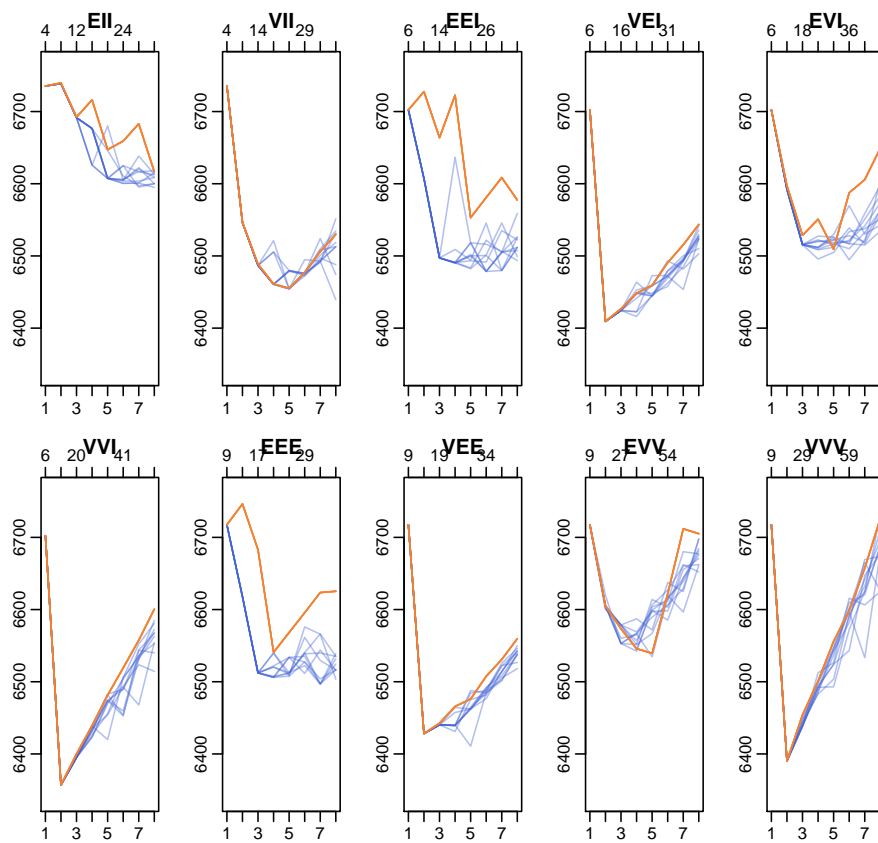
```
> savdir <- file.path(mainsav, "2init")
> filenames <- list.files(savdir, pattern=".rds")
> MW214fn <- grep("MW214", filenames, value="TRUE")
> mclustfiles <- grep("mcl.rds", MW214fn, value=TRUE)
> MW214fn <- grep("mcl.rds", MW214fn, value="TRUE", invert=TRUE)
> claraMW <- grep("clara", MW214fn, value=TRUE)
> mclMW <- grep("mclVVV", MW214fn, value=TRUE)
> clarabic <- massbic(claraMW, savdir)
> mclbic <- massbic(mclMW, savdir)
> mclustbic <- readRDS(file.path(savdir, mclustfiles[1]))
```

here some examples of fitted mixtures


We can see, that, subtracting the obvious hiccups of the small erroneous components, `norMmix` has correctly found the 'intended' distribution. This is remarkable, given the small sample size and difficulty of distribution

```
> compplot(s10mw34bic, m1034)
```

unnamed




```
> plot(MW34)
```



3.5 Nonnormal Mixtures

here 2smi and 2var, maybe others as well.

here 2smi:

```
> savdir <- file.path(mainsav, "2smi")
> filenames <- list.files(savdir, pattern=".rds")
> fnclara <- grep("clara_seed", filenames, value=TRUE)
> fnmclVV <- grep("mclVVV_see", filenames, value=TRUE)
> fnmclus <- grep("__mcl.rds", filenames, value=TRUE)
```


While not very spectacular, the graphs show that even at large parameter counts our algorithm closes in on the same values as `mclust`. At these dimensions it is difficult to compare if these are actually equal, or even similar fits, but going by BIC values, it is at the very least equally viable as a working model.

To illustrate, here are the parameter sizes for this simulation:

	EII	VII	EEI	VEI	EVI	VVI	EEE	VEE	EVV	VVV
1	21	21	40	40	40	40	230	230	230	230
2	42	43	61	62	80	81	251	252	460	461
3	63	65	82	84	120	122	272	274	690	692
4	84	87	103	106	160	163	293	296	920	923


```
> complot(clarabic, mclbic, mclustbic, main="Fit of MW34")
```

Fit of MW34



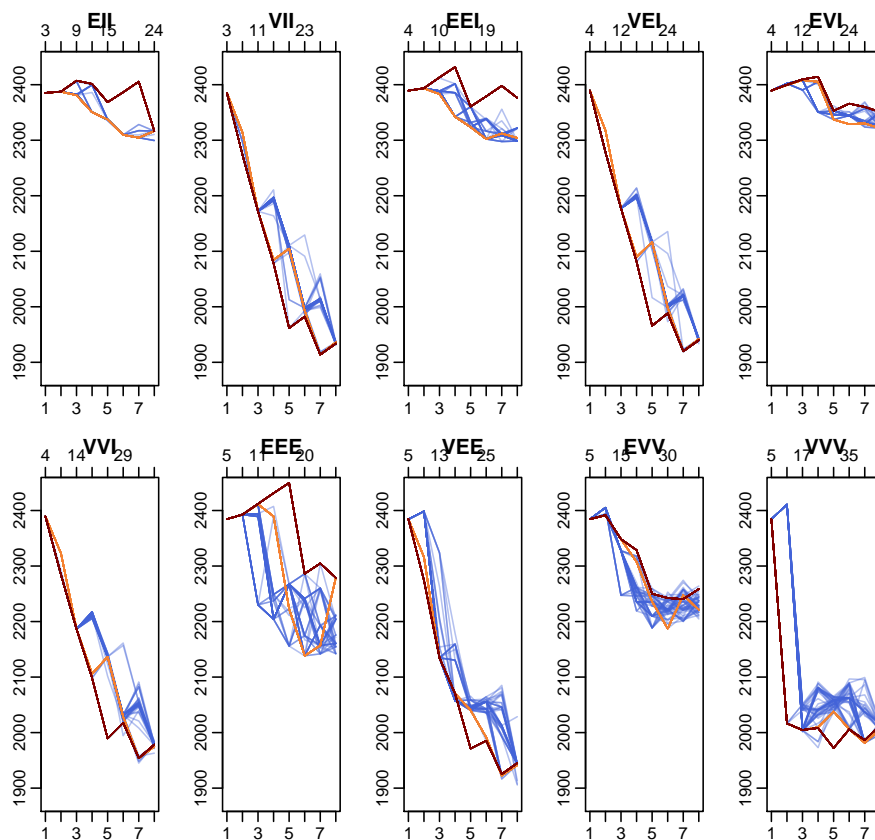
```
5 105 109 124 128 200 204 314 318 1150 1154
6 126 131 145 150 240 245 335 340 1380 1385
7 147 153 166 172 280 286 356 362 1610 1616
8 168 175 187 194 320 327 377 384 1840 1847
```

```
> plot(MW214)
```




```
> compplot(clarabic, mclbic, mclustbic, main="Fit of MW214")
```

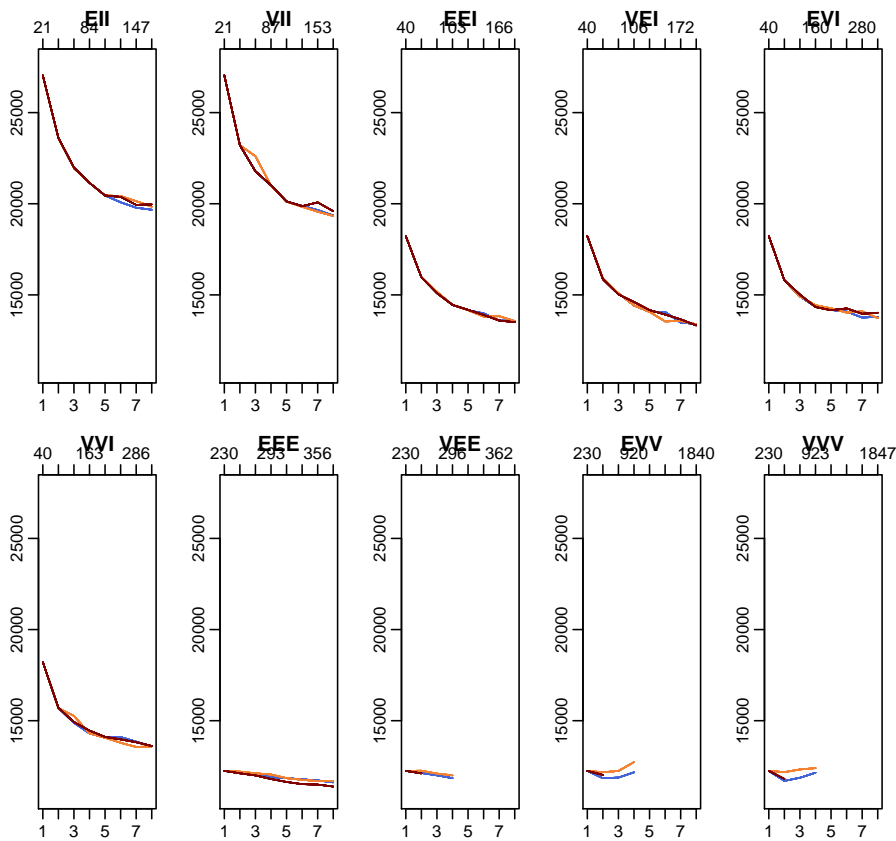
Fit of MW214



```
> f <- readRDS(file.path(savdir, claraMW[28]))  
> ff <- f$fit$nMm[8,8][[1]]  
> plot(ff$norMmix)  
> #points(ff$x)
```



BIC of SMI.12



245 Chapter 4

246 Discussion

247 one shortcoming is time inefficiency. largely due to implementation. mclust has 16'000
248 lines of Fortran code, impossible in the scope of this thesis.

249 proof of concept?? definitely possible to do model selection using a general optimizer.

250 strong points: 'randomness' of clara/optim allows 'confidence intervalls' for selected model
251 flexibility of approach: given an logLik fctn can do mixture fitting w/ arbitrary models

252 further study might include: other presumed component distributions, 'high' dimensions

Bibliography

- Celeux, G. and G. Govaert (1995). Gaussian parsimonious clustering models. *Pattern Recognition* 28(5), 781 – 793.
- Dempster, A. P., N. M. Laird, and D. B. Rubin (1977). Maximum likelihood from incomplete data via the em algorithm. *Journal of the Royal Statistical Society: Series B (Methodological)* 39(1), 1–22.
- L, S., F. M, M. TB, and R. AE. (2016). mclust 5: Clustering, classification and density estimation using gaussian finite mixture models. *R J.* (8(1)), 289–317.
- Marron, J. S. and M. P. Wand (1992). Exact mean integrated squared error. *The Annals of Statistics* 20(2), 712–736.
- McLachlan, G. and D. Peel (2000). *Finite Mixture Models* (1 ed.). Wiley Series in Probability and Statistics. Wiley-Interscience.
- Pearson, K. and O. M. F. E. Henrici (1896). Vii. mathematical contributions to the theory of evolution.—iii. regression, heredity, and panmixia. *Philosophical Transactions of the Royal Society of London. Series A, Containing Papers of a Mathematical or Physical Character* 187, 253–318.

269 Appendix A

270 R Code

271 A.1 llnorMmix

272 Here llnorMmix, since it is the central piece of the package, and 2time.R as an example
273 of a simulation script.

```
#### the llnorMmix function, calculating log likelihood for a given
#### parameter vector
## Author: Nicolas Trutmann 2019-07-06
## Log-likelihood of parameter vector given data
#
# par:    parameter vector
# tx:     transposed sample matrix
# k:      number of components
# model:  assumed distribution model of normal mixture
# trafo:  either centered log ratio or logit
llnorMmix <- function(par, tx, k,
                      trafo=c("clr1", "logit"),
                      model=c("EII","VII","EEI","VEI","EVI",
                              "VVI","EEE","VEE","EVV","VVV")
                      ) {
  stopifnot(is.matrix(tx),
            length(k <- as.integer(k)) == 1, k >= 1)
  p <- nrow(tx)
#   x <- t(x) ## then only needed in (x-mu[,i])^2 i=1..k
# 2. transform
  model <- match.arg(model)
  trafo <- match.arg(trafo)
  l2pi <- log(2*pi)
# 3. calc log-lik
# get w
  w <- if (k==1) 1
    else switch(trafo,
                "clr1" = clr1inv (par[1:(k-1)]),
                "logit"= logitinv(par[1:(k-1)]),
```

```

        stop("invalid 'trafo': ", trafo)
    )
    # start of relevant parameters:
    f <- k + p*k # weights -1 + means +1 => start of alpha
    # get mu
    mu <- matrix(par[k:(f-1L)], p,k)
    f1 <- f      # end of alpha if uniform
    f2 <- f+k-1L # end of alpha if var
    f1.1 <- f1 +1L # start of D. if alpha unif.
    f2.1 <- f1 + k # start of D. if alpha variable
    f11 <- f1 + p-1 # end of D. if D. uniform and alpha uniform
    f12 <- f1 +(p-1)*k # end D. if D. var and alpha unif.
    f21 <- f2 + p-1 # end of D. if D. uniform and alpha variable
    f22 <- f2 +(p-1)*k # end of D. if D. var and alpha var.
    f11.1 <- f11 +1L # start of L if alpha unif D unif
    f21.1 <- f21 +1L # start of L if alpha var D unif
    f12.1 <- f12 +1L # start of L if alpha unif D var
    f22.1 <- f22 +1L # start of L if alpha var D var
    f111 <- f11 + p*(p-1)/2 # end of L if alpha unif D unif
    f211 <- f21 + p*(p-1)/2 # end of L if alpha var D unif
    f121 <- f12 + k*p*(p-1)/2 # end of L if alpha unif D var
    f221 <- f22 + k*p*(p-1)/2 # end of L if alpha var D var
    # initialize f(tx_i) i=1..n vector of density values
    invl <- 0
    # calculate log-lik, see first case for explanation
    switch(model,
    "EII" = {
        alpha <- par[f]
        invalpha <- exp(-alpha)# = 1/exp(alpha)
        for (i in 1:k) {
            rss <- colSums(invalpha*(tx-mu[,i])^2)
            # this is vector of length n=sample size
            # calculates (tx-mu)t * Sigma^-1 * (tx-mu) for diagonal
            # cases.
            invl <- invl+w[i]*exp(-0.5*(p*(alpha+l2pi)+rss))
            # adds likelihood of one component to invl
            # the formula in exp() is the log of likelihood
            # still of length n
        }
    },
    # hereafter differences are difference in dimension in alpha and D.
    # alpha / alpha[i] and D. / D.[,i]
    "VII" = {
        alpha <- par[f:f2]
        for (i in 1:k) {
            rss <- colSums((tx-mu[,i])^2/exp(alpha[i]))
            invl <- invl+w[i]*exp(-0.5*(p*(alpha[i]+l2pi)+rss))
        }
    },

```

```

"EEI" = {
  alpha <- par[f]
  D. <- par[f1.1:f11]
  D. <- c(-sum(D.),D.)
  D. <- D.-sum(D.)/p
  invD <- exp(alpha+D.)
  for (i in 1:k) {
    rss <- colSums((tx-mu[,i])^2/invD)
    invl <- invl+w[i]*exp(-0.5*(p*(alpha+l2pi)+rss))
  }
},
"VEI" = {
  alpha <- par[f:f2]
  D. <- par[f2.1:f21]
  D. <- c(-sum(D.), D.)
  D. <- D.-sum(D.)/p
  for (i in 1:k) {
    rss <- colSums((tx-mu[,i])^2/exp(alpha[i]+D.))
    invl <- invl+w[i]*exp(-0.5*(p*(alpha[i]+l2pi)+rss))
  }
},
"EVI" = {
  alpha <- par[f]
  D. <- matrix(par[f1.1:f12],p-1,k)
  D. <- apply(D.,2, function(j) c(-sum(j), j))
  D. <- apply(D.,2, function(j) j-sum(j)/p)
  for (i in 1:k) {
    rss <- colSums((tx-mu[,i])^2/exp(alpha+D.[,i]))
    invl <- invl+w[i]*exp(-0.5*(p*(alpha+l2pi)+rss))
  }
},
"VVI" = {
  alpha <- par[f:f2]
  D. <- matrix(par[f2.1:f22],p-1,k)
  D. <- apply(D.,2, function(j) c(-sum(j), j))
  D. <- apply(D.,2, function(j) j-sum(j)/p)
  for (i in 1:k) {
    rss <- colSums((tx-mu[,i])^2/exp(alpha[i]+D.[,i]))
    invl <- invl+w[i]*exp(-0.5*(p*(alpha[i]+l2pi)+rss))
  }
},
# here start the non-diagonal cases. main difference is the use
# of backsolve() to calculate tx^t Sigma^-1 tx, works as follows:
# assume Sigma = L D L^t, then Sigma^-1 = (L^t)^-1 D^-1 L^-1
# y = L^-1 tx => tx^t Sigma^-1 tx = y^t D^-1 y
# y = backsolve(L., tx)
"EEE" = {
  alpha <- par[f]
  D. <- par[f1.1:f11]

```

```

D. <- c(-sum(D.), D.)
D. <- D.-sum(D./p)
invD <- exp(alpha+D.)
L. <- diag(1,p)
L.[lower.tri(L., diag=FALSE)] <- par[f11.1:f111]
for (i in 1:k) {
  rss <- colSums(backsolve(L.,(tx-mu[,i]), upper.tri=FALSE)^2/invD)
  invl <- invl+w[i]*exp(-0.5*(p*(alpha+l2pi)+rss))
}
},
"VEE" = {
  alpha <- par[f:f2]
  D. <- par[f2.1:f21]
  D. <- c(-sum(D.), D.)
  D. <- D.-sum(D./p)
  L. <- diag(1,p)
  L.[lower.tri(L., diag=FALSE)] <- par[f21.1:f211]
  for (i in 1:k) {
    rss <- colSums(backsolve(L., (tx-mu[,i]), upper.tri=FALSE)^2/exp(alpha[i]+D.))
    invl <- invl+w[i]*exp(-0.5*(p*(alpha[i]+l2pi)+rss))
  }
},
"EVV" = {
  alpha <- par[f]
  D. <- matrix(par[f1.1:f12],p-1,k)
  D. <- apply(D.,2, function(j) c(-sum(j), j))
  D. <- apply(D.,2, function(j) j-sum(j)/p)
  L.temp <- matrix(par[f12.1:f121],p*(p-1)/2,k)
  for (i in 1:k) {
    L. <- diag(1,p)
    L.[lower.tri(L., diag=FALSE)] <- L.temp[,i]
    rss <- colSums(backsolve(L., (tx-mu[,i]), upper.tri=FALSE)^2/exp(alpha+D.))
    invl <- invl+w[i]*exp(-0.5*(p*(alpha+l2pi)+rss))
  }
},
"VVV" = {
  alpha <- par[f:f2]
  D. <- matrix(par[f2.1:f22],p-1,k)
  D. <- apply(D.,2, function(j) c(-sum(j), j))
  D. <- apply(D.,2, function(j) j-sum(j)/p)
  invalpha <- exp(rep(alpha, each=p)+D.)
  L.temp <- matrix(par[f22.1:f221],p*(p-1)/2,k)
  L. <- diag(1,p)
  for (i in 1:k) {
    L.[lower.tri(L., diag=FALSE)] <- L.temp[,i]
    rss <- colSums(backsolve(L., (tx-mu[,i]), upper.tri=FALSE)^2/invalpha[,i])
    invl <- invl+w[i]*exp(-0.5*(p*(alpha[i]+l2pi)+rss))
  }
},

```

```

    ## otherwise
    stop("invalid model:", model)
  )
  ## return  $\sum_{i=1}^n \log( f(tx_i) )$  :
  sum(log(invl))
}

sllnorMmix <- function(x, obj, trafo=c("clr1", "logit")) {
  stopifnot(is.character(model <- obj$model))
  trafo <- match.arg(trafo)
  llnorMmix(nMm2par(obj, model=model),
            tx = t(x), k = obj$k,
            model=model, trafo=trafo)
}

## log-likelihood function relying on mvtnorm function
#
# par:   parameter vector as calculated by nMm2par
# x:     matrix of samples
# k:     number of cluster
# trafo: transformation of weights
# model: assumed model of the distribution
llmvtnorm <- function(par, x, k,
                     trafo=c("clr1", "logit"),
                     model=c("EII", "VII", "EEI", "VEI", "EVI",
                             "VVI", "EEE", "VEE", "EVV", "VVV"))
  ) {
  stopifnot(is.matrix(x),
            length(k <- as.integer(k)) == 1, k >= 1)
  model <- match.arg(model)
  trafo <- match.arg(trafo)
  p <- ncol(x)
  nmm <- par2nMm(par, p, k, model=model, trafo=trafo)
  ## FIXME (speed!): dmvnorm(*, sigma= S) will do a chol(S) for each component
  ## ----- *instead* we already have LDL' and chol(S) = sqrt(D) L' !!
  ## another par2*() function should give L and D, or from that chol(Sigma), rather than S
  w <- nmm$w
  mu <- nmm$mu
  sig <- nmm$Sigma
  y <- 0
  for (i in 1:k) {
    y <- y + w[i]*mvtnorm::dmvnorm(x, mean=mu[,i], sigma=sig[,i])
  }
  sum(log(y))
}

```

274 A.2 Example Simulation Script

275 here e.g. 2init.R and write some remarks on it.

```
## Intent: analyse time as function of p,k,n
nmmdir <- normalizePath("~/BachelorArbeit/norMmix.Rcheck/")
savdir <- normalizePath("~/BachelorArbeit/Rscripts/2time")
stopifnot(dir.exists(nmmdir), dir.exists(savdir))
library(norMmix, lib.loc=nmmdir)
library(mclust)
## at n=500,p=2 can do about 250xfitnMm(x,1:10) in 24h
seeds <- 1:10
sizes <- c(500, 1000, 2000)
nmm <- list(MW214, MW34, MW51)
## => about 100 cases
# for naming purposes
nmnames <- c("MW214", "MW34", "MW51")
sizenames <- c("500", "1000", "2000")
files <- vector(mode="character")
for (nm in 1:3) {
  for (size in sizes) {
    set.seed(2019); x <- rnorMmix(size, nmm[[nm]])
    for (seed in seeds) {
      set.seed(2019+seed)
      r <- tryCatch(fitnMm(x, k=1:8,
                          optREPORT=1e4, maxit=1e4),
                    error = identity)
      filename <- sprintf("%s_size=%0.4d_seed=%0.2d.rds",
                          nmnames[nm], size, seed)
      files <- append(files, filename)
      cat("===> saving to file:", filename, "\n")
      saveRDS(list(fit=r), file=file.path(savdir, filename))
    }
  }
}
fillis <- list()
for (i in seq_along(sizes)) {
  for (j in seq_along(nmnames)) {
    # for lack of AND matching, OR match everything else and invert
    ret <- grep(paste(sizenames[-i], nmnames[-j], sep="|"),
                files, value=TRUE, invert=TRUE)
    fillis[[paste0(sizenames[i], nmnames[j])]] <- ret
  }
}
epfl(fillis, savdir)
```


276 **Appendix B**

277 **Further Plots**

278 here further plots:

279 **B.1 Ch3**

280 dsfasdf

Declaration of Originality

The signed declaration of originality is a component of every semester paper, Bachelor's thesis, Master's thesis and any other degree paper undertaken during the course of studies, including the respective electronic versions.

Lecturers may also require a declaration of originality for other written papers compiled for their courses.

I hereby confirm that I am the sole author of the written work here enclosed and that I have compiled it in my own words. Parts excepted are corrections of form and content by the supervisor .

Title of work (in block letters):

Authored by (in block letters):

For papers written by groups the names of all authors are required.

Name(s):

First name(s):

Master	Student

With my signature I confirm that

- I have committed none of the forms of plagiarism described in the Citation etiquette information sheet.
- I have documented all methods, data and processes truthfully.
- I have not manipulated any data.
- I have mentioned all persons who were significant facilitators of the work .
- I am aware that the work may be screened electronically for plagiarism.
- I have understood and followed the guidelines in the document *Scientific Works in Mathematics*.

Place, date:

Signature(s):

Zurich August 19th 2009	bla

For papers written by groups the names of all authors are required. Their signatures collectively guarantee the entire content of the written paper.